

# CIND-820 Capstone Project: An ML Tool to Detect Heart Disease

- Robert M. Pineau
- 941-049-371
- Supervisor: Dr. Ceni Babaoglu

## Install Required Modules:

```
1 #!pip install matplotlib
2 #!pip install graphviz
```

## Load Required Libraries:

```
1 import sys
2 from google.colab import drive
3 import math
4 from statistics import mean, stdev
5 import pandas as pd
6 import numpy as np
7 from scipy import stats
8 import plotly
9 import matplotlib.pyplot as plt
10 import seaborn as sns
11
12 from sklearn import tree
13 from sklearn.naive_bayes import MultinomialNB
14 from sklearn.naive_bayes import GaussianNB
15 from sklearn.linear_model import LogisticRegression
16 from sklearn.svm import SVC
17 from sklearn.tree import DecisionTreeClassifier
18 from sklearn.neighbors import KNeighborsClassifier
19 from sklearn.ensemble import RandomForestClassifier
20 from xgboost import XGBClassifier
21
22 from sklearn.metrics import confusion_matrix
23 from sklearn.model_selection import train_test_split
24 from sklearn.metrics import classification_report
25 from sklearn import metrics
26 from sklearn.metrics import accuracy_score
27 from sklearn.metrics import f1_score
28
29 from sklearn.model_selection import KFold
30 from sklearn.model_selection import cross_validate
```

```

31 from sklearn.model_selection import cross_val_score
32 from sklearn.model_selection import RandomizedSearchCV, GridSearchCV
33 import imblearn
34 from imblearn.over_sampling import SMOTE
35
36 import graphviz

```

## Obtain Data-Set from Google Drive:

```

1 # Mounting google colab, this will prompt first time each session.
2 drive.mount('/content/drive',force_remount=True)
3 dataset_file = "/content/drive/My Drive/Colab Notebooks/heart_statlog_cleveland_hungary-fi
4 df=pd.read_csv(dataset_file,sep=',')
5 df.name = "Original Data-Set"
6 print(df.name)
7 df.head(3)

```

Mounted at /content/drive  
Original Data-Set

	age	sex	chest pain type	resting bp s	cholesterol	fasting blood sugar	resting ecg	max heart rate	exercise angina	oldpeak	s1
0	40	1	2	140	289	0	0	172	0	0.0	
1	49	0	3	160	180	0	0	156	0	1.0	

## Clean up Column Names:

```

1 #Rename the columns to be nicer, no spaces.
2 df=df.rename(columns={"age": "Age", "sex": "Sex", "chest pain type": "ChestPainType", "res
3 df=df.rename(columns={"cholesterol":"Cholesterol","fasting blood sugar": "FastingBloodSuga
4 df=df.rename(columns={"chest pain type": "ChestPainType", "resting bp s": "RestingBP_s", "
5 df=df.rename(columns={"resting ecg": "RestingECG", "max heart rate": "MaxHeartRate", "exer
6 df=df.rename(columns={"oldpeak":"OldPeak", "ST slope": "ST_Slope", "target": "Target"})
7 df.name = "Original Data-Set"
8 df.head(3)

```

	Age	Sex	ChestPainType	RestingBP_s	Cholesterol	FastingBloodSugar	RestingECG	Major vessels (Ma
0	40	1	2	140	289	0	0	
1	49	0	3	160	180	0	0	
2	37	1	2	130	283	0	1	

## Datatypes and Quantities:

```
1 #Check data types.
2 print(df.dtypes)
```

```
Age                int64
Sex                int64
ChestPainType      int64
RestingBP_s        int64
Cholesterol         int64
FastingBloodSugar  int64
RestingECG         int64
MaxHeartRate       int64
ExerciseAngina     int64
OldPeak            float64
ST_Slope           int64
Target             int64
dtype: object
```

```
1 #Datatypes, counts, etc.
2 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1190 entries, 0 to 1189
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age                   1190 non-null  int64
1   Sex                   1190 non-null  int64
2   ChestPainType         1190 non-null  int64
3   RestingBP_s           1190 non-null  int64
4   Cholesterol            1190 non-null  int64
5   FastingBloodSugar     1190 non-null  int64
6   RestingECG            1190 non-null  int64
7   MaxHeartRate          1190 non-null  int64
8   ExerciseAngina        1190 non-null  int64
9   OldPeak               1190 non-null  float64
10  ST_Slope              1190 non-null  int64
11  Target                1190 non-null  int64
dtypes: float64(1), int64(11)
memory usage: 111.7 KB
```

### Check for Missing and NULL entries:

```
1 #Check for NULL or missing entries. (none)
2 print(df.isna().any())
3 print("\n\n")
4 print(df.isnull().any())
```

```
Age                False
Sex                False
ChestPainType      False
RestingBP_s        False
```

```

Cholesterol      False
FastingBloodSugar  False
RestingECG       False
MaxHeartRate     False
ExerciseAngina   False
OldPeak          False
ST_Slope         False
Target           False
dtype: bool

```

```

Age              False
Sex              False
ChestPainType    False
RestingBP_s      False
Cholesterol      False
FastingBloodSugar  False
RestingECG       False
MaxHeartRate     False
ExerciseAngina   False
OldPeak          False
ST_Slope         False
Target           False
dtype: bool

```

## Check for Duplicate Entries:

```

1 #Look for rows that are 100% identical to each other.
2 #
3 dup_count = sum(df.duplicated())
4 print(f"There are {dup_count} duplicate rows in this dataset.\n")
5
6 #Dropping any duplicate entries.
7 df = df.drop_duplicates(ignore_index = True)
8 df.name = "Original Data-Set"
9
10 df.info()

```

There are 272 duplicate rows in this dataset.

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 918 entries, 0 to 917
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age                   918 non-null   int64
1   Sex                   918 non-null   int64
2   ChestPainType         918 non-null   int64
3   RestingBP_s           918 non-null   int64
4   Cholesterol            918 non-null   int64
5   FastingBloodSugar     918 non-null   int64
6   RestingECG            918 non-null   int64

```

```

7   MaxHeartRate      918 non-null    int64
8   ExerciseAngina    918 non-null    int64
9   OldPeak           918 non-null    float64
10  ST_Slope           918 non-null    int64
11  Target             918 non-null    int64
dtypes: float64(1), int64(11)
memory usage: 86.2 KB

```

## Check for Out-Of-Bound Entries for Nominal and Binary attributes:

```

1 #Check for out of bound entries(outliers) for nominal and binary attributes (including the
2 #Since all nominal and binary attributes have a valid contiguous integer range, ie 0-1, or
3 #we only need to look for those outside the range.
4 valid_values = {'Sex': [0,1], 'ChestPainType': [1,2,3,4], 'FastingBloodSugar': [0,1], 'Res
5
6 for col in ('Sex', 'ChestPainType', 'FastingBloodSugar', 'RestingECG', 'ExerciseAngina', '
7   valid = np.array(valid_values[col])
8   max_valid = valid.max()
9   min_valid = valid.min()
10  print(f"For attribute '{col}': Valid MAX: {max_valid}, Valid MIN: {min_valid}")
11  these_outliers = df[((df[col] < min_valid) | (df[col] > max_valid))]
12
13  if (these_outliers.shape[0] > 1):
14      print(f"For attribute '{col}': There are {these_outliers.shape[0]} outliers:\n")
15      print(these_outliers)
16      print("\n\n")
17  elif (these_outliers.shape[0] == 1):
18      print(f"For attribute '{col}': There is {these_outliers.shape[0]} outlier:\n")
19      print(these_outliers)
20      print("\n\n")
21  else:
22      print(f"For attribute '{col}': There are no outliers.\n")

```

```

For attribute 'Sex': Valid MAX: 1, Valid MIN: 0
For attribute 'Sex': There are no outliers.

```

```

For attribute 'ChestPainType': Valid MAX: 4, Valid MIN: 1
For attribute 'ChestPainType': There are no outliers.

```

```

For attribute 'FastingBloodSugar': Valid MAX: 1, Valid MIN: 0
For attribute 'FastingBloodSugar': There are no outliers.

```

```

For attribute 'RestingECG': Valid MAX: 2, Valid MIN: 0
For attribute 'RestingECG': There are no outliers.

```

```

For attribute 'ExerciseAngina': Valid MAX: 1, Valid MIN: 0
For attribute 'ExerciseAngina': There are no outliers.

```

```

For attribute 'ST_Slope': Valid MAX: 3, Valid MIN: 1
For attribute 'ST_Slope': There is 1 outlier:

```

```

Age  Sex  ChestPainType  RestingBP_s  Cholesterol  FastingBloodSugar  \

```

516	68	1	3	150	195	1
	RestingECG	MaxHeartRate	ExerciseAngina	OldPeak	ST_Slope	Target
516	0	132	0	0.0	0	1

For attribute 'Target': Valid MAX: 1, Valid MIN: 0  
 For attribute 'Target': There are no outliers.

## Remove Out-Of-Bound Entry:

```

1 #From above, there is a problem with one entry regarding the ST_Slope attribute, it is zer
2 #
3 #The documentation at https://ieee-dataport.org/open-access/heart-disease-dataset-comprehe
4 #Shows a range of 0-2, but in the definition of the mapped nominal values it shows:
5 #
6 # -- Value 1: upsloping
7 # -- Value 2: flat
8 # -- Value 3: downsloping
9 #
10 print(df['ST_Slope'].value_counts().sort_index())
11 print("\n")
12 #
13 #
14 #Since there is only one entry out of range, making the assumption that the correct range
15 #
16 #Will simply drop this entry.
17 df = df[df['ST_Slope'] != 0]
18 df.name = "Original Data-Set"
19
20 df.info()

```

```

0      1
1    395
2    459
3     63
Name: ST_Slope, dtype: int64

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 917 entries, 0 to 917
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age                   917 non-null   int64
1   Sex                   917 non-null   int64
2   ChestPainType         917 non-null   int64
3   RestingBP_s           917 non-null   int64
4   Cholesterol            917 non-null   int64
5   FastingBloodSugar     917 non-null   int64

```

```

6   RestingECG          917 non-null    int64
7   MaxHeartRate        917 non-null    int64
8   ExerciseAngina      917 non-null    int64
9   OldPeak             917 non-null    float64
10  ST_Slope            917 non-null    int64
11  Target              917 non-null    int64
dtypes: float64(1), int64(11)
memory usage: 93.1 KB

```

## Basic Statistics for All Attributes:

```

1 #Basic Statistics of the dataset.(Measures of Center/Central Tendency, and Measures of Var
2 df.describe().T

```

	count	mean	std	min	25%	50%	75%	max
<b>Age</b>	917.0	53.495093	9.425601	28.0	47.0	54.0	60.0	77.0
<b>Sex</b>	917.0	0.789531	0.407864	0.0	1.0	1.0	1.0	1.0
<b>ChestPainType</b>	917.0	3.251908	0.931502	1.0	3.0	4.0	4.0	4.0
<b>RestingBP_s</b>	917.0	132.377317	18.515114	0.0	120.0	130.0	140.0	200.0
<b>Cholesterol</b>	917.0	198.803708	109.443764	0.0	173.0	223.0	267.0	603.0
<b>FastingBloodSugar</b>	917.0	0.232279	0.422517	0.0	0.0	0.0	0.0	1.0
<b>RestingECG</b>	917.0	0.604144	0.806161	0.0	0.0	0.0	1.0	2.0
<b>MaxHeartRate</b>	917.0	136.814613	25.473732	60.0	120.0	138.0	156.0	202.0
<b>ExerciseAngina</b>	917.0	0.404580	0.491078	0.0	0.0	0.0	1.0	1.0
<b>OldPeak</b>	917.0	0.888332	1.066749	-2.6	0.0	0.6	1.5	6.2
<b>ST_Slope</b>	917.0	1.637950	0.607270	1.0	1.0	2.0	2.0	3.0
<b>Target</b>	917.0	0.552890	0.497466	0.0	0.0	1.0	1.0	1.0

## Visualizing All Attributes:

```

1 #Assigning descriptive labels for all possible values for all nominal/binary attributes.
2 #
3 labels = {'Sex': ['Female', 'Male'], 'ChestPainType': ['Typical Angina', 'A-Typical Angina',
4             'FastingBloodSugar': ['<= 120 mg/dl', '> 120 mg/dl'], 'RestingECG': ['Normal', '
5             'ExerciseAngina':['No', 'Yes'], 'ST_Slope': ['Upsloping', 'Flat', 'Downsloping']
6
7

```

```


1 #Creating subsets of the data for a series of interesting plots to help with visualization

```

```
2 #
3 #Breaking up the dataset into two groups, those with heart disease and those without.
4 with_heart_disease = df[df['Target'] == 1]
5 no_heart_disease = df[df['Target'] == 0]
6
7 #Breaking up the dataset into four groups, by Sex, and those with heart disease and those
8 with_heart_disease_male = df[(df['Target'] == 1) & (df['Sex'] == 1)]
9 with_heart_disease_female = df[(df['Target'] == 1) & (df['Sex'] == 0)]
10 no_heart_disease_male = df[(df['Target'] == 0) & (df['Sex'] == 1)]
11 no_heart_disease_female = df[(df['Target'] == 0) & (df['Sex'] == 0)]
12
13 #For these groups remove the "Sex" column from the data.
14 with_heart_disease_male = with_heart_disease_male.drop("Sex",axis=1)
15 with_heart_disease_female = with_heart_disease_female.drop("Sex",axis=1)
16 no_heart_disease_male = no_heart_disease_male.drop("Sex",axis=1)
17 no_heart_disease_female = no_heart_disease_female.drop("Sex",axis=1)

1 #Visualizing the distribution of the categorical attributes, including the class variable.
2 ax=1
3 plt.figure(figsize=(15,15))
4
5 for col in ('Sex', 'ChestPainType', 'FastingBloodSugar', 'RestingECG', 'ExerciseAngina','S
6     plt.subplot(4,2,ax)
7     these_labels = labels[col]
8     plt.title(f"Distribution for attribute '{col}':")
9     plt.pie(df[col].value_counts().sort_index(),
10         autopct = '%1.1f%%', labels=these_labels,
11         explode=tuple([0.1] * len(these_labels)))
12     plt.axis('equal')
13     ax+=1
14
15 plt.suptitle('Nominal/Binary Attribute Distribution',y=1.01, size = 16, color = 'black', w
16 plt.tight_layout()
17 plt.savefig("nominal_dist.pdf",dpi=1200, bbox_inches='tight')
```



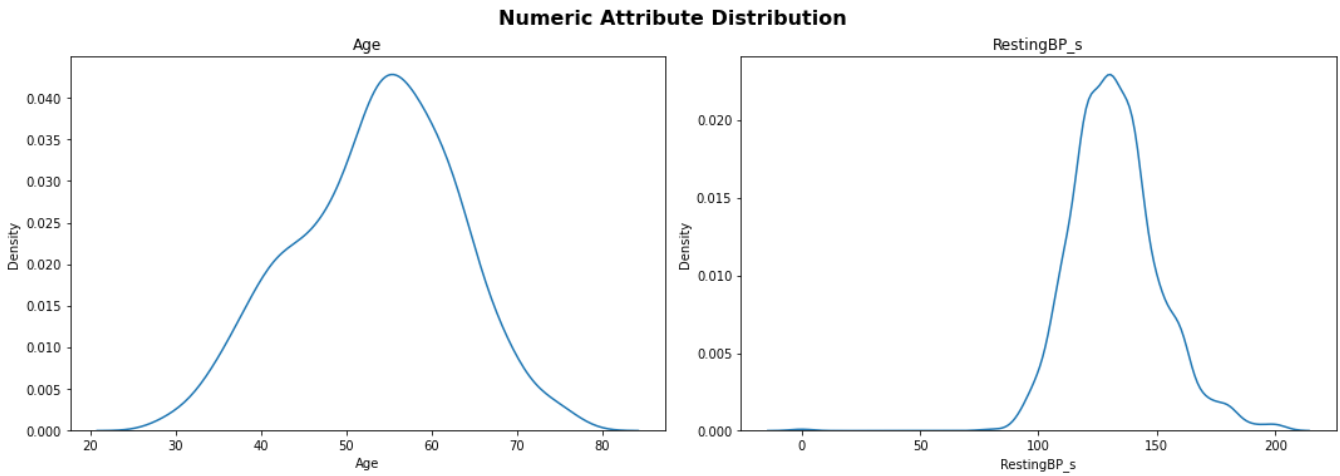
**Nominal/Binary Attribute Distribution**

```

1 #Visualizing the overall distribution of the numeric attributes.
2 plt.figure(figsize=(15,15))
3
4 ax=1
5 for col in ('Age', 'RestingBP_s', 'Cholesterol', 'MaxHeartRate', 'OldPeak'):
6     plt.subplot(3,2,ax)
7     plt.title(col)
8     sns.kdeplot(x=df[col])

```

```
9     ax += 1
10
11 plt.suptitle('Numeric Attribute Distribution',y=1.01, size = 16, color = 'black', weight='
12 plt.tight_layout()
13 plt.savefig("numeric_dist.pdf",dpi=1200, bbox_inches='tight')
```



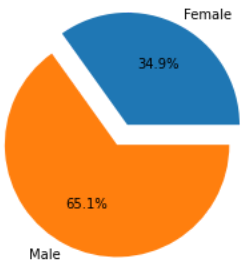
```

1 #Visualizing the distribution of the categorical attributes, by target
2 ax=1
3 plt.figure(figsize=(15,20))
4 plt.axis('equal')
5
6 for col in ('Sex','ChestPainType', 'FastingBloodSugar', 'RestingECG', 'ExerciseAngina','ST
7     plt.subplot(6,2,ax)
8     these_labels = labels[col]
9     plt.title(f"Distribution for attribute '{col}' Without Heart Disease:")
10    plt.pie(no_heart_disease[col].value_counts().sort_index(),
11        autopct = '%1.1f%%', labels=these_labels,
12        explode=tuple([0.1] * len(these_labels)))
13    ax+=1
14    plt.subplot(6,2,ax)
15    these_labels = labels[col]
16    plt.title(f"Distribution for attribute '{col}' With Heart Disease:")
17    plt.pie(with_heart_disease[col].value_counts().sort_index(),
18        autopct = '%1.1f%%', labels=these_labels,
19        explode=tuple([0.1] * len(these_labels)))
20    ax+=1
21
22 plt.suptitle('Nominal/Binary Attribute Distribution by Target',y=1.01, size = 16, color =
23 plt.tight_layout()
24 plt.savefig("nominal_dist_by_target.pdf",dpi=1200, bbox_inches='tight')

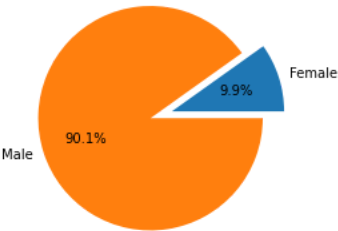
```

Nominal/Binary Attribute Distribution by Target

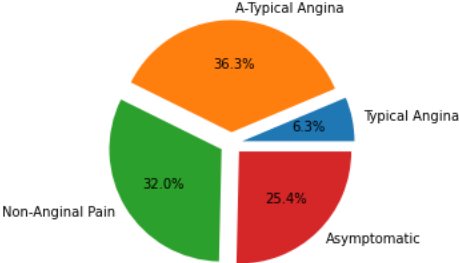
Distribution for attribute 'Sex' Without Heart Disease:




Distribution for attribute 'Sex' With Heart Disease:




Distribution for attribute 'ChestPainType' Without Heart Disease:




Distribution for attribute 'ChestPainType' With Heart Disease:




Distribution for attribute 'FastingBloodSugar' Without Heart Disease:




Distribution for attribute 'FastingBloodSugar' With Heart Disease:




Distribution for attribute 'RestingECG' Without Heart Disease:




Distribution for attribute 'RestingECG' With Heart Disease:



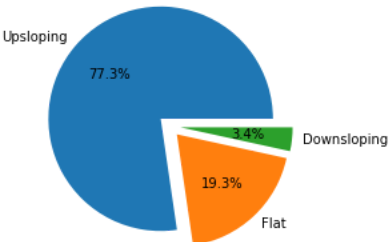
Distribution for attribute 'ExerciseAngina' Without Heart Disease:




Distribution for attribute 'ExerciseAngina' With Heart Disease:



Distribution for attribute 'ST\_Slope' Without Heart Disease:

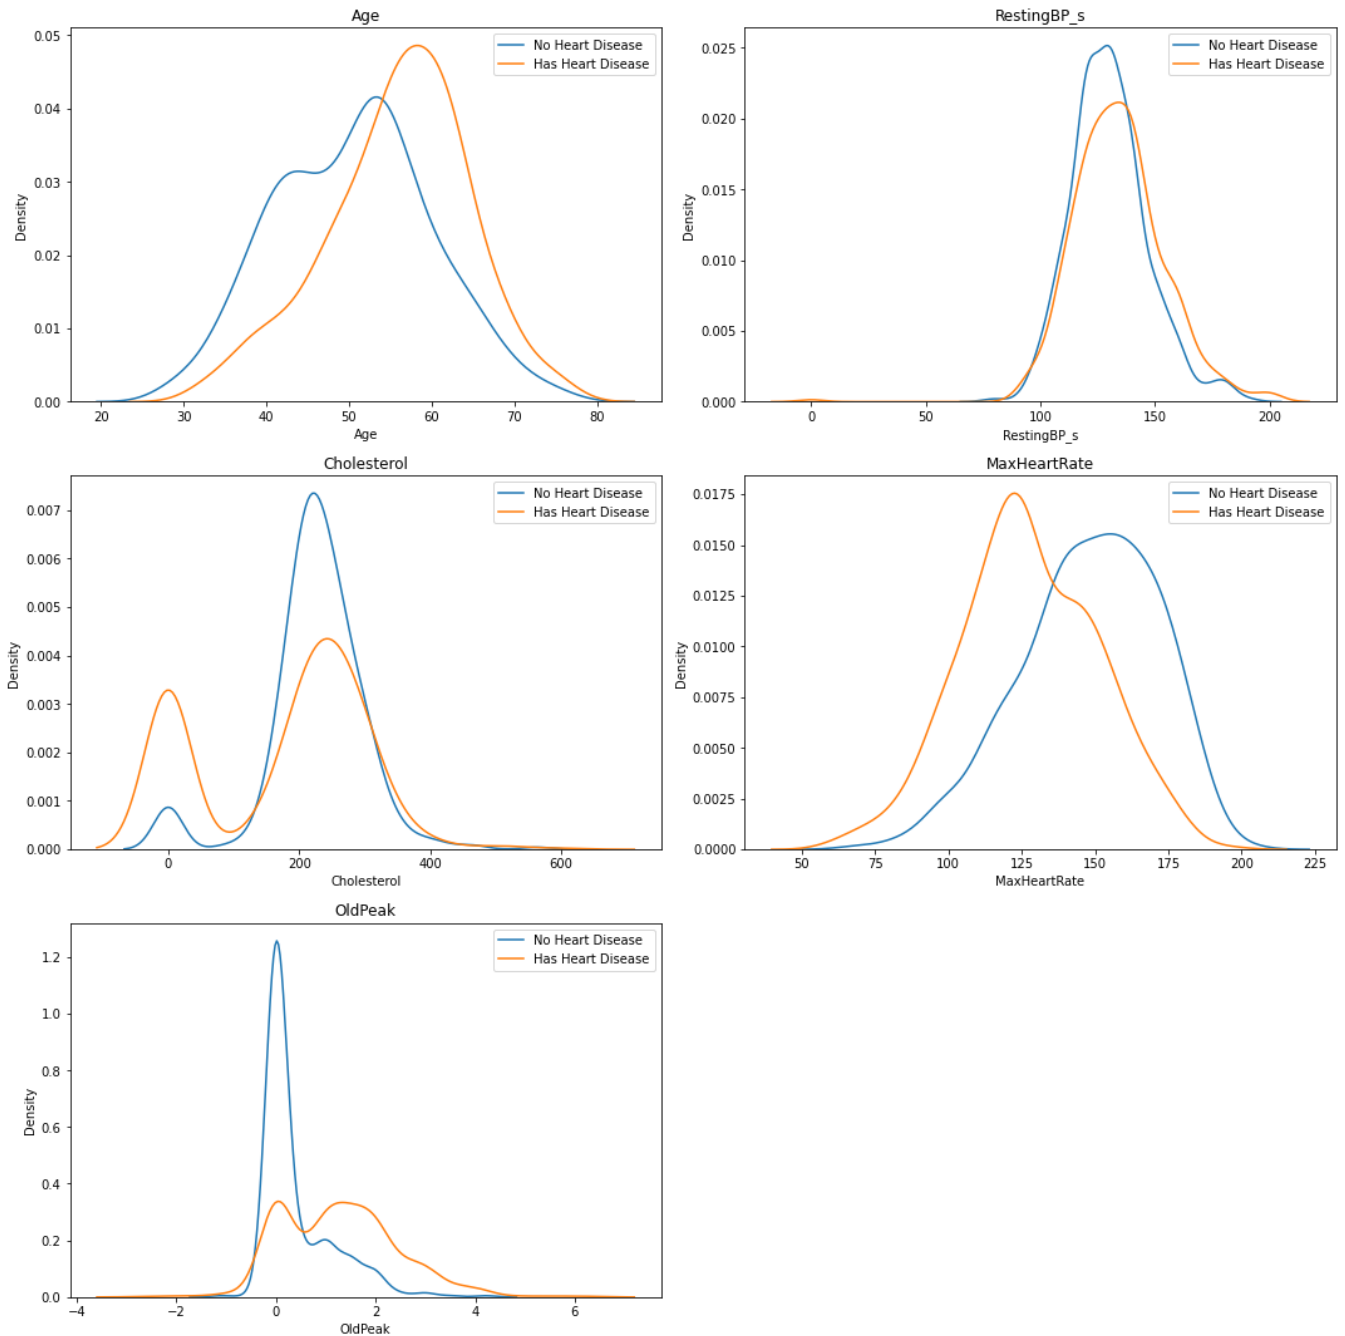


Distribution for attribute 'ST\_Slope' With Heart Disease:



```
1 #Visualizing the distribution of the numeric attributes by Target:
2 plt.figure(figsize=(15,15))
3
4 ax=1
5 for col in ('Age', 'RestingBP_s', 'Cholesterol', 'MaxHeartRate', 'OldPeak'):
6     plt.subplot(3,2,ax)
7     plt.title(col)
8     sns.kdeplot(x=no_heart_disease[col],label = "No Heart Disease")
9     sns.kdeplot(x=with_heart_disease[col],label = "Has Heart Disease")
10    plt.legend()
11    ax += 1
12
13 plt.suptitle('Numeric Attribute Distribution by Target',y=1.01, size = 16, color = 'black')
14 plt.tight_layout()
15 plt.savefig("numeric_dist_by_target.pdf",dpi=1200, bbox_inches='tight')
```

## Numeric Attribute Distribution by Target






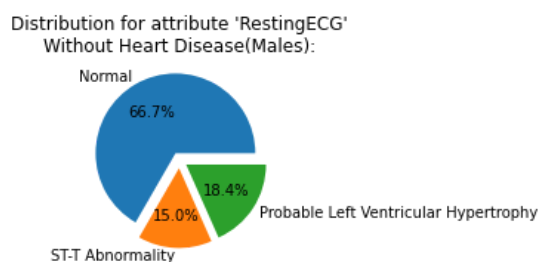

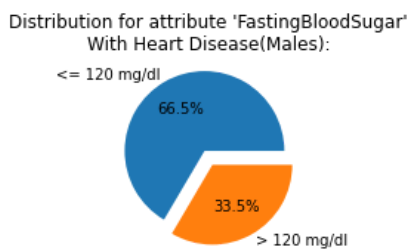

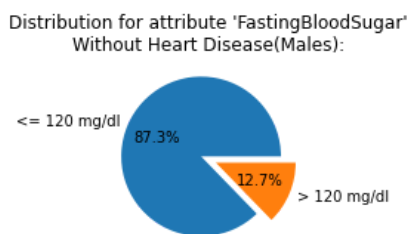
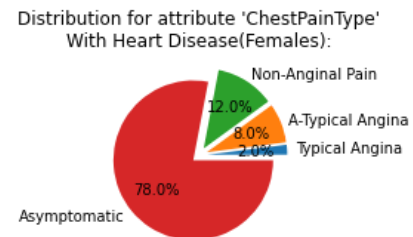
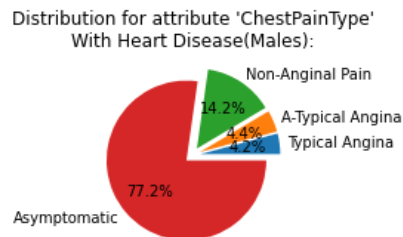
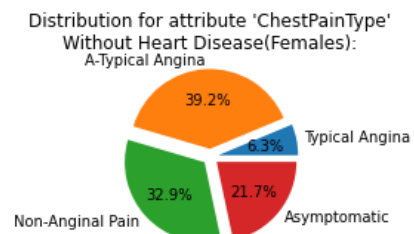
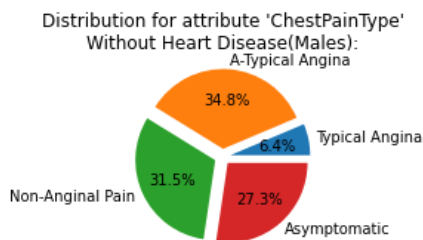
```

1 #Visualizing the distribution of the categorical attributes, by target and by sex
2 #For report purposes, breaking this up into two separate pages.
3 #This one for attributes 'ChestPainType', 'FastingBloodSugar', 'RestingECG'
4 #
5 ax=1
6 plt.figure(figsize=(15,30))
7 plt.axis('equal')
8
9 for col in ('ChestPainType', 'FastingBloodSugar', 'RestingECG'):
10     plt.subplot(12,2,ax)
11     these_labels = labels[col]
12     plt.title(f"Distribution for attribute '{col}'\nWithout Heart Disease(Males):")
13     plt.pie(no_heart_disease_male[col].value_counts().sort_index(),
14             autopct = '%1.1f%%', labels=these_labels,

```

```
15     explode=tuple([0.1] * len(these_labels)))
16     ax+=1
17
18     plt.subplot(12,2,ax)
19     these_labels = labels[col]
20     plt.title(f"Distribution for attribute '{col}'\nWithout Heart Disease(Females):")
21     plt.pie(no_heart_disease_female[col].value_counts().sort_index(),
22             autopct = '%1.1f%%', labels=these_labels,
23             explode=tuple([0.1] * len(these_labels)))
24     ax+=1
25
26     plt.subplot(12,2,ax)
27     these_labels = labels[col]
28     plt.title(f"Distribution for attribute '{col}'\nWith Heart Disease(Males):")
29     plt.pie(with_heart_disease_male[col].value_counts().sort_index(),
30             autopct = '%1.1f%%', labels=these_labels,
31             explode=tuple([0.1] * len(these_labels)))
32     ax+=1
33
34     plt.subplot(12,2,ax)
35     these_labels = labels[col]
36     plt.title(f"Distribution for attribute '{col}'\nWith Heart Disease(Females):")
37     plt.pie(with_heart_disease_female[col].value_counts().sort_index(),
38             autopct = '%1.1f%%', labels=these_labels,
39             explode=tuple([0.1] * len(these_labels)))
40     ax+=1
41
42 plt.suptitle('Nominal/Binary Attribute Distribution by Target and by Sex',y=1.01, size = 1
43 plt.tight_layout()
44 plt.savefig("nominal_dist_by_target_by_sex1.pdf",dpi=1200, bbox_inches='tight')
```

### Nominal/Binary Attribute Distribution by Target and by Sex



```

1 #Visualizing the distribution of the categorical attributes, by target and by sex
2 #For report purposes, breaking this up into two separate pages.
3 #This one for attributes 'ExerciseAngina','ST_Slope'
4 #
5 ax=1
6 plt.figure(figsize=(15,30))
7 plt.axis('equal')
8
9 for col in ('ExerciseAngina','ST_Slope'):
10     plt.subplot(12,2,ax)
11     these_labels = labels[col]
12     plt.title(f"Distribution for attribute '{col}'\nWithout Heart Disease(Males):")

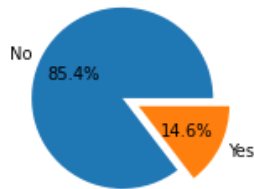
```




```
13 plt.pie(no_heart_disease_male[col].value_counts().sort_index(),
14         autopct = '%1.1f%%', labels=these_labels,
15         explode=tuple([0.1] * len(these_labels)))
16 ax+=1
17
18 plt.subplot(12,2,ax)
19 these_labels = labels[col]
20 plt.title(f"Distribution for attribute '{col}'\nWithout Heart Disease(Females):")
21 plt.pie(no_heart_disease_female[col].value_counts().sort_index(),
22         autopct = '%1.1f%%', labels=these_labels,
23         explode=tuple([0.1] * len(these_labels)))
24 ax+=1
25
26 plt.subplot(12,2,ax)
27 these_labels = labels[col]
28 plt.title(f"Distribution for attribute '{col}'\nWith Heart Disease(Males):")
29 plt.pie(with_heart_disease_male[col].value_counts().sort_index(),
30         autopct = '%1.1f%%', labels=these_labels,
31         explode=tuple([0.1] * len(these_labels)))
32 ax+=1
33
34 plt.subplot(12,2,ax)
35 these_labels = labels[col]
36 plt.title(f"Distribution for attribute '{col}'\nWith Heart Disease(Females):")
37 plt.pie(with_heart_disease_female[col].value_counts().sort_index(),
38         autopct = '%1.1f%%', labels=these_labels,
39         explode=tuple([0.1] * len(these_labels)))
40 ax+=1
41
42 plt.suptitle('Nominal/Binary Attribute Distribution by Target and by Sex, Cont\'d',y=1.01,
43 plt.tight_layout()
44 plt.savefig("nominal_dist_by_target_by_sex2.pdf",dpi=1200, bbox_inches='tight')
```

## Nominal/Binary Attribute Distribution by Target and by Sex, Cont'd


Distribution for attribute 'ExerciseAngina'  
Without Heart Disease(Males):




Distribution for attribute 'ExerciseAngina'  
Without Heart Disease(Females):




Distribution for attribute 'ExerciseAngina'  
With Heart Disease(Males):




Distribution for attribute 'ExerciseAngina'  
With Heart Disease(Females):



Distribution for attribute 'ST\_Slope'  
Without Heart Disease(Males):



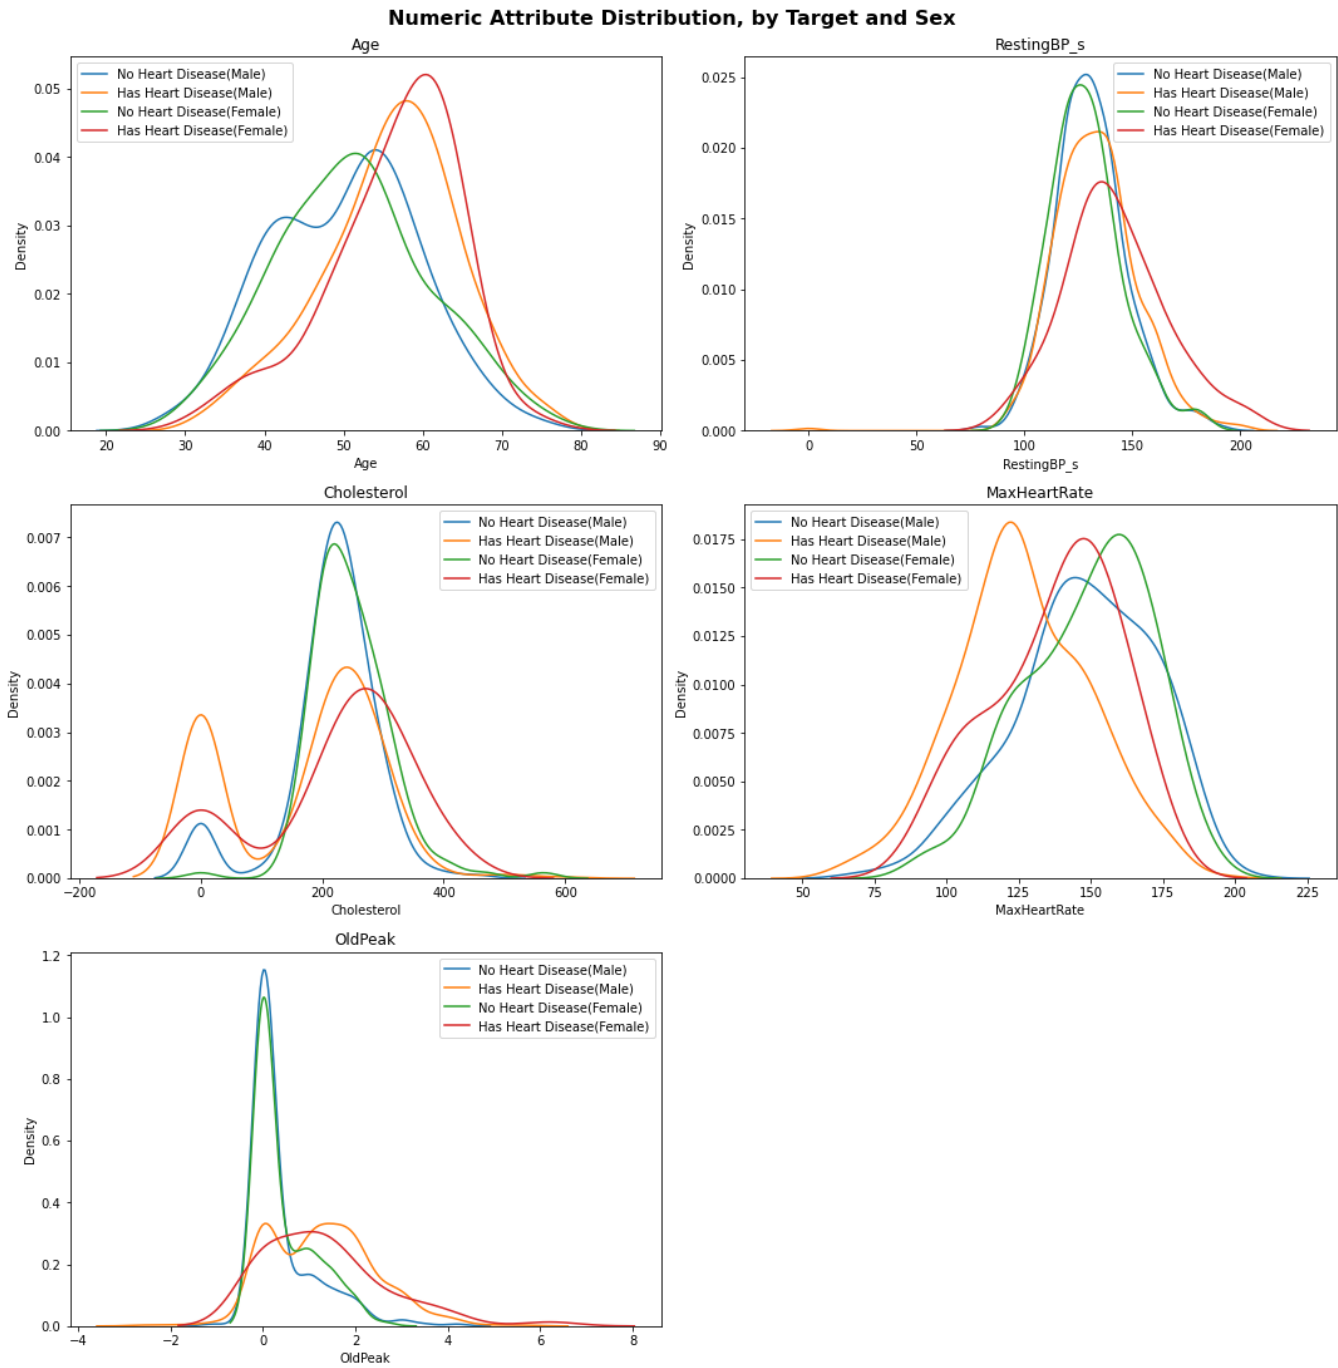
Distribution for attribute 'ST\_Slope'  
Without Heart Disease(Females):



```

1 #Visualizing the distribution of the numerical attributes, by target and by sex
2 plt.figure(figsize=(15,15))
3
4 ax=1
5 for col in ('Age', 'RestingBP_s', 'Cholesterol', 'MaxHeartRate', 'OldPeak'):
6     plt.subplot(3,2,ax)
7     plt.title(col)
8     sns.kdeplot(x=no_heart_disease_male[col],label = "No Heart Disease(Male)")
9     sns.kdeplot(x=with_heart_disease_male[col],label = "Has Heart Disease(Male)")
10    sns.kdeplot(x=no_heart_disease_female[col],label = "No Heart Disease(Female)")
11    sns.kdeplot(x=with_heart_disease_female[col],label = "Has Heart Disease(Female)")
12    plt.legend()
13    ax += 1
14
15 plt.suptitle('Numeric Attribute Distribution, by Target and Sex',y=1.01, size = 16, color
16 plt.tight_layout()
17 plt.savefig("numeric_dist_by_target_by_sex.pdf",dpi=1200, bbox_inches='tight')

```



**Outlier Detection:**

```
1 #Check for outliers on numeric attributes
2 #Using for outlier detection three methods.
3 #Note: for the next stage in this project, Module 3,
4 #one or more of these outlier detection methods will be used.
5 #For now, we only want to see how many outliers per attribute are detected with each appro
6 #
7 #Methods:
8 # #1 1.5IQR range
9 # #2 mean +/- 3*ST-DEV (same as GT Absolute(Z-Score))
10 # #3 Rejecting those with a value of zero (based on visualization, only needed for 'Chole
11
12 def IQR1_5_upper(data, col):
13     Q3 = np.quantile(data[col], 0.75)
14     Q1 = np.quantile(data[col], 0.25)
15     IQR = Q3 - Q1
16     return(Q3+(1.5*IQR))
17
18 def IQR1_5_lower(data, col):
19     Q3 = np.quantile(data[col], 0.75)
20     Q1 = np.quantile(data[col], 0.25)
21     IQR = Q3 - Q1
22     return(Q1-(1.5*IQR))
23
24
25 for col in ('Age', 'RestingBP_s', 'Cholesterol', 'MaxHeartRate', 'OldPeak'):
26     upper1 = IQR1_5_upper(df,col)
27     lower1 = IQR1_5_lower(df,col)
28     stdev3 = 3*df[col].std()
29     mean = df[col].mean()
30     upper2 = mean + stdev3
31     lower2 = mean - stdev3
32
33     these_outliers1 = df[(df[col] < lower1) | (df[col] > upper1)]
34     these_outliers2 = df[(df[col] < lower2) | (df[col] > upper2)]
35     these_outliers3 = df[df[col] == 0]
36
37     print(f"For attribute '{col}': The mean is {mean}, stdev3 is {stdev3}")
38     print(f"For 1.5IQR the lower range is {lower1} the upper range is {upper1}")
39     print(f"For mean +/- 3STDEV the lower range is {lower2} the upper range is {upper2}")
40     print(f"\n")
41
42     print(f"Using 1.5IQR Method:")
43     if (these_outliers1.shape[0] > 1):
44         print(f"For attribute '{col}': There are {these_outliers1.shape[0]} outliers:\n")
45         print(these_outliers1)
46         print("\n")
47     elif (these_outliers1.shape[0] == 1):
48         print(f"For attribute '{col}': There is {these_outliers1.shape[0]} outlier:\n")
49         print(these_outliers1)
50         print("\n")
51     else:
```

```
52     print(f"For attribute '{col}': There are no outliers.\n")
53     print("\n")
54
55     print(f"Using mean +/- 3STDEV Method:")
56     if (these_outliers2.shape[0] > 1):
57         print(f"For attribute '{col}': There are {these_outliers2.shape[0]} outliers:\n")
58         print(these_outliers2)
59         print("\n")
60     elif (these_outliers2.shape[0] == 1):
61         print(f"For attribute '{col}': There is {these_outliers2.shape[0]} outlier:\n")
62         print(these_outliers2)
63         print("\n")
64     else:
65         print(f"For attribute '{col}': There are no outliers.\n")
66         print("\n")
67
68     if(col == 'Cholesterol'):
69         print(f"Identifying 'zero' values(for 'Cholesterol') Method:")
70         if (these_outliers3.shape[0] > 1):
71             print(f"For attribute '{col}': There are {these_outliers3.shape[0]} outliers:\n")
72             print(these_outliers3)
73             print("\n")
74         elif (these_outliers3.shape[0] == 1):
75             print(f"For attribute '{col}': There is {these_outliers3.shape[0]} outlier:\n")
76             print(these_outliers3)
77             print("\n")
78         else:
79             print(f"For attribute '{col}': There are no outliers.\n")
80             print("\n")
81
82 print("\n\n")
```

For attribute 'Age': The mean is 53.49509269356598, stdev3 is 28.276802628148687  
For 1.5IQR the lower range is 27.5 the upper range is 79.5  
For mean +/- 3STDEV the lower range is 25.218290065417293 the upper range is 81.77189

Using 1.5IQR Method:  
For attribute 'Age': There are no outliers.

Using mean +/- 3STDEV Method:  
For attribute 'Age': There are no outliers.

For attribute 'RestingBP\_s': The mean is 132.3773173391494, stdev3 is 55.545341581484  
For 1.5IQR the lower range is 90.0 the upper range is 170.0  
For mean +/- 3STDEV the lower range is 76.83197575766476 the upper range is 187.92265

Using 1.5IQR Method:

For attribute 'RestingBP\_s': There are 28 outliers:

	Age	Sex	ChestPainType	RestingBP_s	Cholesterol	FastingBloodSugar	\
109	39	1	2	190	241	0	
123	58	0	2	180	393	0	
189	53	1	4	180	285	0	
190	46	1	4	180	280	0	
241	54	1	4	200	198	0	
274	45	0	2	180	295	0	
275	59	1	3	180	213	0	
278	57	0	4	180	347	0	
314	53	1	4	80	0	0	
365	64	0	4	200	0	0	
372	63	1	4	185	0	0	
399	61	1	3	200	0	1	
411	54	1	4	180	0	1	
423	60	1	3	180	0	0	
449	55	1	3	0	0	0	
475	59	1	4	178	0	1	
550	55	1	4	172	260	0	
585	57	1	2	180	285	1	
592	61	1	4	190	287	1	
673	59	0	4	174	249	0	
702	59	1	1	178	270	0	
725	55	0	4	180	327	0	
732	56	0	4	200	288	1	
759	54	1	2	192	283	0	
774	66	0	4	178	228	1	
780	64	0	4	180	325	0	
855	68	1	3	180	274	1	
880	52	1	3	172	199	1	

	RestingECG	MaxHeartRate	ExerciseAngina	OldPeak	ST_Slope	Target
109	0	106	0	0.0	1	0
123	0	110	1	1.0	2	1
189	1	120	1	1.5	2	1

## Data Manipulation, for Outliers, and Model Considerations:

```

1 #Before we can run multinomial Naive Bayes we must remove any negative numbers in the data
2 mins = df.min()
3 print(mins)
4 print("\n\n")
5 #There are only negative values for attribute 'OldPeak'
6 #Applying a simple shift to eliminate any negatives.
7
8 df_no_negs = df.copy()
9 df_no_negs.name = "No Negatives Data-Set"
10
11 df_no_negs['OldPeak'] = df_no_negs['OldPeak'] + abs(df_no_negs['OldPeak'].min())
12

```

```


13 #Visualizing the overall distribution of 'OldPeak' before and after modification for negat
14 plt.figure(figsize=(10,5))
15
16 plt.subplot(1,2,1)
17 plt.title('OldPeak')
18 sns.kdeplot(x=df['OldPeak'])
19
20 plt.subplot(1,2,2)
21 plt.title('OldPeakNoNeg')
22 sns.kdeplot(x=df_no_negs['OldPeak'])
23
24 plt.suptitle('Visualizing the Transformation for OldPeak',y=1.01, size = 16, color = 'blac
25 plt.tight_layout()
26 plt.savefig("oldpeak_transformation.pdf",dpi=1200, bbox_inches='tight')
27

```

```

Age                28.0
Sex                0.0
ChestPainType      1.0
RestingBP_s        0.0
Cholesterol         0.0
FastingBloodSugar  0.0
RestingECG         0.0
MaxHeartRate       60.0
ExerciseAngina     0.0
OldPeak            -2.6
ST_Slope           1.0
Target             0.0
dtype: float64


```



```

1 #Only addressing outliers for attribute Cholesterol, specifically the instances where Chol
2 #All other outliers appear in much smaller quantities.
3
4 df_outliers_addressed = df_no_negs.copy()
5 df_outliers_addressed.name = "Outliers Addressed Data-Set"
6
7 cholesterol_mean = df_outliers_addressed['Cholesterol'][df_outliers_addressed['Cholesterol
8
9 df_outliers_addressed['Cholesterol'].replace(to_replace=0.0, value=cholesterol_mean, inpla
10
11
12 #Visualizing the overall distribution Cholesterol before and after dealing with outliers.
13 plt.figure(figsize=(10,5))
14
15 plt.subplot(1,2,1)
16 plt.title('Cholesterol')
17 sns.kdeplot(x=df['Cholesterol'])
18
19 plt.subplot(1,2,2)
20 plt.title('Cholesterol(No Outliers)')
21 sns.kdeplot(x=df_outliers_addressed['Cholesterol'])
22
23 plt.suptitle('Visualizing the Transformation for Cholesterol',y=1.01, size = 16, color = '
24 plt.tight_layout()

```



### Normalize Numeric Data for potential model training:

```

1 #Normalizing numeric data to see if this helps, or hinders the accuracy of the ML models.
2 df_normalized = df_outliers_addressed.copy()
3 df_normalized.name = "Normalized Data-Set"

```



```

4
5 for col in ('Age', 'RestingBP_s', 'Cholesterol', 'MaxHeartRate', 'OldPeak'):
6     df_normalized[col] = (df[col]-df[col].min())/(df[col].max()-df[col].min())

```

### Create ONE-HOT columns for all categorical attributes:

```

1 #For Nominal & Binary (ie categorical) attributes, perform one-hot conversion.
2 #convert only categorical variables/features to dummy/one-hot features
3 cat_cols = ['Sex','ChestPainType', 'FastingBloodSugar', 'RestingECG', 'ExerciseAngina','ST
4
5 df_onehot = pd.get_dummies(df, columns=cat_cols, prefix = cat_cols)
6 df_onehot.name = "Original Data-Set, After ONEHOT"
7
8 df_onehot_no_negs = pd.get_dummies(df_no_negs, columns=cat_cols, prefix = cat_cols)
9 df_onehot_no_negs.name = "No negatives Data-Set, After ONEHOT"
10
11 df_onehot_outliers_addressed = pd.get_dummies(df_outliers_addressed, columns=cat_cols, pre
12 df_onehot_outliers_addressed.name = "Outliers Addressed Data-Set, After ONEHOT"
13
14 df_onehot_normalized = pd.get_dummies(df_normalized, columns=cat_cols, prefix = cat_cols)
15 df_onehot_normalized.name = "Normalized Data-Set, After ONEHOT"
16

```

### Use SMOTE to Balance the Class Variable:

```


1 # Balance the dataset using SMOTE.
2
3 def do_SMOTE(df, classCol="Target"):
4     oversample = SMOTE()
5     X = df.drop(columns=classCol)
6     Y = df[classCol]
7     X, Y = oversample.fit_resample(X, Y)
8     X[classCol] = Y
9     X.name = df.name+", After SMOTE"
10    return(X)
11
12
13
14 df_smote = do_SMOTE(df)
15 df_no_negs_smote = do_SMOTE(df_no_negs)
16
17 df_normalized_smote = do_SMOTE(df_normalized)
18 df_outliers_addressed_smote = do_SMOTE(df_outliers_addressed)
19
20 df_onehot_smote = do_SMOTE(df_onehot)
21 df_onehot_no_negs_smote = do_SMOTE(df_onehot_no_negs)
22 df_onehot_outliers_addressed_smote = do_SMOTE(df_onehot_outliers_addressed)

```

```

23 df_onehot_normalized_smote = do_SMOTE(df_onehot_normalized)
24
25
26 #Visualizing the distribution of the class variable before and after SMOTE for the basic d
27 ax=1
28 plt.figure(figsize=(10,6))
29
30 for this_df in (df,df_smote):
31     plt.subplot(1,2,ax)
32     these_labels = labels['Target']
33     plt.title(f"{this_df.name}:")
34     plt.pie(this_df['Target'].value_counts().sort_index(),
35             autopct = '%1.1f%%', labels=these_labels,
36             explode=tuple([0.1] * len(these_labels)))
37     plt.axis('equal')
38     ax+=1
39
40 plt.suptitle('Distribution for Class Variable Before and After SMOTE',y=1.01, size = 16, c
41 plt.tight_layout()
42 plt.savefig("smote.pdf",dpi=1200, bbox_inches='tight')
43

```



## ML Algorithms:

```

1 def do_DT(df,levels,class_col_name,verbose=0):
2     #not disabling randomness.

```

```
3  #np.random.seed(0)
4
5  # Split dataset into training set and test set
6  feature_names=df.columns[df.columns != class_col_name ]
7  # 80% training and 20% test
8  X_train, X_test, Y_train, Y_test = train_test_split(df.loc[:, feature_names], df[class_c
9
10  clf = tree.DecisionTreeClassifier(max_depth=levels,criterion='gini')
11  clf = clf.fit(X_train, Y_train)
12  if (verbose >= 1):
13      print(f"Successfully trained the decision tree for {levels} levels...")
14
15  # Let's make the prdictions on the test set  that we set aside earlier using the trained
16  Y_pred = clf.predict(X_test)
17
18  cf=confusion_matrix(Y_test, Y_pred)
19  tn, fp, fn, tp=cf.ravel()
20  tpr=0.0
21  fpr=0.0
22  tpr = tp/(tp+fp)
23  fpr = fp/(fp+tn)
24  fnr = fn/(fn+tp)
25
26  if (verbose >= 2):
27      print ("Confusion Matrix")
28      print(cf)
29      print("")
30      print ("TP: ", tp,", FP: ", fp,", TN: ", tn,", FN:", fn)
31      print ("TPR: ",tpr,", FPR: ",fpr, "FNR: ",fnr)
32
33  #print precision, recall, and accuracy from the perspective of each of the class (0 and
34  if (verbose >= 2):
35      print(classification_report(Y_test, Y_pred, digits=3))
36
37  accuracy = accuracy_score(Y_test, Y_pred)
38  f1_weighted = f1_score(Y_test, Y_pred,average='weighted')
39
40  if (verbose >= 1):
41      print(f"Accuracy is: {accuracy}")
42      print(f"F1 Weighted is: {f1_weighted}")
43      print("")
44
45  return(accuracy,f1_weighted,tpr,fpr,fnr)


1 def do_mnNB(df,class_col_name,verbose=0):
2     #not disabling randomness.
3     #np.random.seed(0)
4
5     # Split dataset into training set and test set
6     feature_names=df.columns[df.columns != class_col_name ]
```

```
7 # 80% training and 20% test
8 X_train, X_test, Y_train, Y_test = train_test_split(df.loc[:, feature_names], df[class_c
9
10 #Create a MultiNomial NB Classifier
11 nb = MultinomialNB()
12
13 #Train the model using the training sets
14 nb.fit(X_train, Y_train)
15
16 #Predict the response for test dataset
17 Y_pred = nb.predict(X_test)
18
19 if (verbose >= 2):
20     print ("Total Columns (including class)",len(df.columns))
21     print("Classes ",nb.classes_)
22     print("Number of records for classes ",nb.class_count_)
23     print("Log prior probability for classes ", nb.class_log_prior_)
24     print("Log conditional probability for each feature given a class\n",nb.feature_log_pr
25
26 cf=confusion_matrix(Y_test, Y_pred)
27 tn, fp, fn, tp=cf.ravel()
28 tpr = tp/(tp+fp)
29 fpr = fp/(fp+tn)
30 fnr = fn/(fn+tp)
31
32 if (verbose >= 2):
33     print ("Confusion Matrix")
34     print(cf)
35     print("")
36     print ("TP: ", tp,", FP: ", fp,", TN: ", tn,", FN:", fn)
37     print ("TPR: ",tpr,", FPR: ",fpr, "FNR: ",fnr)
38
39 if (verbose >= 2):
40     print(classification_report(Y_test, Y_pred, digits=3))
41
42 accuracy = accuracy_score(Y_test, Y_pred)
43 f1_weighted = f1_score(Y_test, Y_pred,average='weighted')
44
45 if (verbose >= 1):
46     print(f"Accuracy is: {accuracy}")
47     print(f"F1 Weighted is: {f1_weighted}")
48     print("")
49
50 return(accuracy,f1_weighted,tpr,fpr,fnr)

1 def do_gaNB(df,class_col_name,verbose=0):
2     #not disabling randomness.
3     #np.random.seed(0)
4
5     # Split dataset into training set and test set
```

```
6 feature_names=df.columns[df.columns != class_col_name ]
7 # 80% training and 20% test
8 X_train, X_test, Y_train, Y_test = train_test_split(df.loc[:, feature_names], df[class_c
9
10 #Create a Gaussian NB Classifier
11 nb = GaussianNB()
12
13 #Train the model using the training sets
14 nb.fit(X_train, Y_train)
15
16 #Predict the response for test dataset
17 Y_pred = nb.predict(X_test)
18
19 if (verbose >= 2):
20     print ("Total Columns (including class)",len(df.columns))
21     print("Number of records for classes ",nb.class_count_)
22
23 cf=confusion_matrix(Y_test, Y_pred)
24 tn, fp, fn, tp=cf.ravel()
25 tpr = tp/(tp+fp)
26 fpr = fp/(fp+tn)
27 fnr = fn/(fn+tp)
28
29 if (verbose >= 2):
30     print ("Confusion Matrix")
31     print(cf)
32     print("")
33     print ("TP: ", tp,", FP: ", fp,", TN: ", tn,", FN:", fn)
34     print ("TPR: ",tpr,", FPR: ",fpr, "FNR: ",fnr)
35
36 if (verbose >= 2):
37     print(classification_report(Y_test, Y_pred, digits=3))
38
39 accuracy = accuracy_score(Y_test, Y_pred)
40 f1_weighted = f1_score(Y_test, Y_pred,average='weighted')
41
42 if (verbose >= 1):
43     print(f"Accuracy is: {accuracy}")
44     print(f"F1 Weighted is: {f1_weighted}")
45     print("")
46
47 return(accuracy,f1_weighted,tpr,fpr,fnr)

1 def do_LR(df,class_col_name,verbose=0):
2     #not disabling randomness.
3     #np.random.seed(0)
4
5     # Split dataset into training set and test set
6     feature_names=df.columns[df.columns != class_col_name ]
7     # 80% training and 20% test
```

```
8 X_train, X_test, Y_train, Y_test = train_test_split(df.loc[:, feature_names], df[class_c
9
10 lr = LogisticRegression(max_iter=2000)
11
12 #Train the model using the training sets
13 lr.fit(X_train, Y_train)
14
15 #Predict the response for test dataset
16 Y_pred = lr.predict(X_test)
17
18 if (verbose >= 2):
19     print ("Total Columns (including class)",len(df.columns))
20
21 cf=confusion_matrix(Y_test, Y_pred)
22 tn, fp, fn, tp=cf.ravel()
23 tpr = tp/(tp+fp)
24 fpr = fp/(fp+tn)
25 fnr = fn/(fn+tp)
26
27 if (verbose >= 2):
28     print ("Confusion Matrix")
29     print(cf)
30     print("")
31     print ("TP: ", tp,",", FP: ", fp,",", TN: ", tn,",", FN:",", fn)
32     print ("TPR: ",tpr,",", FPR: ",fpr,",", FNR: ",fnr)
33
34 if (verbose >= 2):
35     print(classification_report(Y_test, Y_pred, digits=3))
36
37 accuracy = accuracy_score(Y_test, Y_pred)
38 f1_weighted = f1_score(Y_test, Y_pred,average='weighted')
39
40 if (verbose >= 1):
41     print(f"Accuracy is: {accuracy}")
42     print(f"F1 Weighted is: {f1_weighted}")
43     print("")
44
45 return(accuracy,f1_weighted,tpr,fpr,fnr)

1 def do_KNN(df,class_col_name,verbose=0):
2     #not disabling randomness.
3     #np.random.seed(0)
4
5     # Split dataset into training set and test set
6     feature_names=df.columns[df.columns != class_col_name ]
7     # 80% training and 20% test
8     X_train, X_test, Y_train, Y_test = train_test_split(df.loc[:, feature_names], df[class_c
9
10     knn = KNeighborsClassifier()
11
```

```
12 #Train the model using the training sets
13 knn.fit(X_train, Y_train)
14
15 #Predict the response for test dataset
16 Y_pred = knn.predict(X_test)
17
18 if (verbose >= 2):
19     print ("Total Columns (including class)",len(df.columns))
20
21 cf=confusion_matrix(Y_test, Y_pred)
22 tn, fp, fn, tp=cf.ravel()
23 tpr = tp/(tp+fp)
24 fpr = fp/(fp+tn)
25 fnr = fn/(fn+tp)
26
27 if (verbose >= 2):
28     print ("Confusion Matrix")
29     print(cf)
30     print("")
31     print ("TP: ", tp,", FP: ", fp,", TN: ", tn,", FN:", fn)
32     print ("TPR: ",tpr,", FPR: ",fpr, "FNR: ",fnr)
33
34 if (verbose >= 2):
35     print(classification_report(Y_test, Y_pred, digits=3))
36
37 accuracy = accuracy_score(Y_test, Y_pred)
38 f1_weighted = f1_score(Y_test, Y_pred,average='weighted')
39
40 if (verbose >= 1):
41     print(f"Accuracy is: {accuracy}")
42     print(f"F1 Weighted is: {f1_weighted}")
43     print("")
44
45 return(accuracy,f1_weighted,tpr,fpr,fnr)
```

```
1 def do_RF(df,class_col_name,verbose=0):
2     #not disabling randomness.
3     #np.random.seed(0)
4
5     # Split dataset into training set and test set
6     feature_names=df.columns[df.columns != class_col_name ]
7     # 80% training and 20% test
8     X_train, X_test, Y_train, Y_test = train_test_split(df.loc[:, feature_names], df[class_c
9
10     rf = RandomForestClassifier()
11
12     #Train the model using the training sets
13     rf.fit(X_train, Y_train)
14
15     #Predict the response for test dataset
```

```
16 Y_pred = rf.predict(X_test)
17
18 if (verbose >= 2):
19     print ("Total Columns (including class)",len(df.columns))
20
21 cf=confusion_matrix(Y_test, Y_pred)
22 tn, fp, fn, tp=cf.ravel()
23 tpr = tp/(tp+fp)
24 fpr = fp/(fp+tn)
25 fnr = fn/(fn+tp)
26
27 if (verbose >= 2):
28     print ("Confusion Matrix")
29     print(cf)
30     print("")
31     print ("TP: ", tp,", FP: ", fp,", TN: ", tn,", FN:", fn)
32     print ("TPR: ",tpr,", FPR: ",fpr, "FNR: ",fnr)
33
34 if (verbose >= 2):
35     print(classification_report(Y_test, Y_pred, digits=3))
36
37 accuracy = accuracy_score(Y_test, Y_pred)
38 f1_weighted = f1_score(Y_test, Y_pred,average='weighted')
39
40 if (verbose >= 1):
41     print(f"Accuracy is: {accuracy}")
42     print(f"F1 Weighted is: {f1_weighted}")
43     print("")
44
45 return(accuracy,f1_weighted,tpr,fpr,fnr)


1 def do_SVM(df,class_col_name,verbose=0):
2     #not disabling randomness.
3     #np.random.seed(0)
4
5     # Split dataset into training set and test set
6     feature_names=df.columns[df.columns != class_col_name ]
7     # 80% training and 20% test
8     X_train, X_test, Y_train, Y_test = train_test_split(df.loc[:, feature_names], df[class_c
9
10     svm = SVC()
11
12     #Train the model using the training sets
13     svm.fit(X_train, Y_train)
14
15     #Predict the response for test dataset
16     Y_pred = svm.predict(X_test)
17
18     if (verbose >= 2):
19         print ("Total Columns (including class)",len(df.columns))
```



```
20
21 cf=confusion_matrix(Y_test, Y_pred)
22 tn, fp, fn, tp=cf.ravel()
23 tpr = tp/(tp+fp)
24 fpr = fp/(fp+tn)
25 fnr = fn/(fn+tp)
26
27
28 if (verbose >= 2):
29     print ("Confusion Matrix")
30     print(cf)
31     print("")
32     print ("TP: ", tp,", FP: ", fp,", TN: ", tn,", FN:", fn)
33     print ("TPR: ",tpr,", FPR: ",fpr, "FNR: ",fnr)
34
35 if (verbose >= 2):
36     print(classification_report(Y_test, Y_pred, digits=3))
37
38 accuracy = accuracy_score(Y_test, Y_pred)
39 f1_weighted = f1_score(Y_test, Y_pred,average='weighted')
40
41 if (verbose >= 1):
42     print(f"Accuracy is: {accuracy}")
43     print(f"F1 Weighted is: {f1_weighted}")
44     print("")
45
46 return(accuracy,f1_weighted,tpr,fpr,fnr)


1 def do_XGB(df,class_col_name,verbose=0):
2     #not disabling randomness.
3     #np.random.seed(0)
4
5     # Split dataset into training set and test set
6     feature_names=df.columns[df.columns != class_col_name ]
7     # 80% training and 20% test
8     X_train, X_test, Y_train, Y_test = train_test_split(df.loc[:, feature_names], df[class_c
9
10    xgb = XGBClassifier()
11
12    #Train the model using the training sets
13    xgb.fit(X_train, Y_train)
14
15    #Predict the response for test dataset
16    Y_pred = xgb.predict(X_test)
17
18    if (verbose >= 2):
19        print ("Total Columns (including class)",len(df.columns))
20
21    cf=confusion_matrix(Y_test, Y_pred)
22    tn, fp, fn, tp=cf.ravel()
```

```
23 tpr = tp/(tp+fp)
24 fpr = fp/(fp+tn)
25 fnr = fn/(fn+tp)
26
27 if (verbose >= 2):
28     print ("Confusion Matrix")
29     print(cf)
30     print("")
31     print ("TP: ", tp,", FP: ", fp,", TN: ", tn,", FN:", fn)
32     print ("TPR: ",tpr,", FPR: ",fpr, "FNR: ",fnr)
33
34 if (verbose >= 2):
35     print(classification_report(Y_test, Y_pred, digits=3))
36
37 accuracy = accuracy_score(Y_test, Y_pred)
38 f1_weighted = f1_score(Y_test, Y_pred,average='weighted')
39
40 if (verbose >= 1):
41     print(f"Accuracy is: {accuracy}")
42     print(f"F1 Weighted is: {f1_weighted}")
43     print("")
44
45 return(accuracy,f1_weighted,tpr,fpr,fnr)
```

## Initial Run of All ML Algorithms:

```
1 #Initial Run of all ML Algorithms just to make sure everything works correctly.
2 #Using original data:
3
4 for i in range(3,11):
5     print(f"DT with {i} levels:")
6     do_DT(df,i,'Target',5)
7
8 print(f"MN_NB:")
9 do_mnNB(df_no_negs,'Target',5)
10
11 print(f"GA_NB:")
12 do_gaNB(df,'Target',5)
13
14 print(f"LR:")
15 do_LR(df,'Target',5)
16
17 print(f"KNN:")
18 do_KNN(df,'Target',5)
19
20 print(f"RF:")
21 do_RF(df,'Target',5)
22
23 print(f"SVM:")
```

```

24 do_SVM(df, 'Target',5)
25
26 print(f"XGB:")
27 do_XGB(df, 'Target',5)

```

DT with 3 levels:

Successfully trained the decision tree for 3 levels...

Confusion Matrix

```

[[76 15]
 [13 80]]

```

TP: 80 , FP: 15 , TN: 76 , FN: 13

TPR: 0.8421052631578947 , FPR: 0.16483516483516483 FNR: 0.13978494623655913

	precision	recall	f1-score	support
0	0.854	0.835	0.844	91
1	0.842	0.860	0.851	93
accuracy			0.848	184
macro avg	0.848	0.848	0.848	184
weighted avg	0.848	0.848	0.848	184

Accuracy is: 0.8478260869565217

F1 Weighted is: 0.8477901120361805

DT with 4 levels:

Successfully trained the decision tree for 4 levels...

Confusion Matrix

```

[[71 13]
 [11 89]]

```

TP: 89 , FP: 13 , TN: 71 , FN: 11

TPR: 0.8725490196078431 , FPR: 0.15476190476190477 FNR: 0.11

	precision	recall	f1-score	support
0	0.866	0.845	0.855	84
1	0.873	0.890	0.881	100
accuracy			0.870	184
macro avg	0.869	0.868	0.868	184
weighted avg	0.869	0.870	0.869	184

Accuracy is: 0.8695652173913043

F1 Weighted is: 0.8694251824344299

DT with 5 levels:

Successfully trained the decision tree for 5 levels...

Confusion Matrix

```

[[65 16]
 [15 88]]

```

TP: 88 , FP: 16 , TN: 65 , FN: 15

TPR: 0.8461538461538461 , FPR: 0.19753086419753085 FNR: 0.14563106796116504

	precision	recall	f1-score	support
0	0.812	0.802	0.807	81

1	0.846	0.854	0.850	103
accuracy			0.832	184
macro avg	0.829	0.828	0.829	184
weighted avg	0.831	0.832	0.831	184

Accuracy is: 0.8315217391304348

## Validation:

```

1 #Create a subroutine to invoke cross_validate for a given model, and dataset.
2 #
3 def doCV(test_num, test_name, model, this_df, folds=5, classCol='Target', verbose=0):
4     X = this_df.drop(classCol,axis=1)
5     Y = this_df[classCol]
6
7     scoring = {'acc': 'accuracy',
8               'rec': 'recall_macro'}
9     scores = cross_validate(model, X, Y, scoring=scoring,
10                             cv=folds, return_train_score=True)
11
12     ACC = scores['test_acc']
13     ACC_mean = ACC.mean() * 100
14     ACC_std = ACC.std()
15
16     TPR = scores['test_rec']
17     TPR_mean = TPR.mean() * 100
18     TPR_std = TPR.std()
19
20     FNR = 1-TPR
21     FNR_mean = FNR.mean() * 100
22     FNR_std = FNR.std()
23
24     if verbose > 0:
25         print(f"For TEST #{test_num}: {test_name:35s} After Cross-Val, Using Data-Set: {this_d
26
27     #Store all results in an easy lookup-table
28     results[f"LAST_TEST"] = test_num
29     results[f"TEST{test_num}#NAME"] = test_name
30     results[f"TEST{test_num}#DFNAME"] = this_df.name
31     results[f"TEST{test_num}#DF_ROWS"] = this_df.shape[0]
32     results[f"TEST{test_num}#ACCURACY_MEAN"] = ACC_mean
33     results[f"TEST{test_num}#ACCURACY_STD"] = ACC_std
34     results[f"TEST{test_num}#ACCURACY"] = ACC
35     results[f"TEST{test_num}#TPR_MEAN"] = TPR_mean
36     results[f"TEST{test_num}#TPR_STD"] = TPR_std
37     results[f"TEST{test_num}#FNR_MEAN"] = FNR_mean
38     results[f"TEST{test_num}#FNR_STD"] = FNR_std
39     results[f"TEST{test_num}#FNR"] = FNR
40
41 #Create subroutines to display test results, and obtain accuracy and FNR for a given test.

```

```

42 def getTestAccuracy(x):
43     value = results[f"TEST{x}#ACCURACY_MEAN"]
44     return(value)
45
46 def getTestF1(x):
47     value = results[f"TEST{x}#F1_WEIGHTED_MEAN"]
48     return(value)
49
50 def getTestFNR(x):
51     value = results[f"TEST{x}#FNR_MEAN"]
52     return(value)
53
54 def getTestAccVec(x):
55     value = results[f"TEST{x}#ACCURACY"]
56     return(value)
57
58 def getTestFNRVec(x):
59     value = results[f"TEST{x}#FNR"]
60     return(value)
61
62 def getTestDFRows(x):
63     value = results[f"TEST{x}#DF_ROWS"]
64     return(value)
65
66
67 def displayResult(x):
68     name = results[f"TEST{x}#NAME"]
69     df_name = results[f"TEST{x}#DFNAME"]
70     accuracy = results[f"TEST{x}#ACCURACY_MEAN"]
71     #f1_weighted = results[f"TEST{x}#F1_WEIGHTED_MEAN"]
72     tpr = results[f"TEST{x}#TPR_MEAN"]
73     #fpr = results[f"TEST{x}#FPR_MEAN"]
74     fnr = results[f"TEST{x}#FNR_MEAN"]
75
76     print(f"Using ML Model: {name} and {df_name}:")
77     #print(f"TEST #{x}: Average Accuracy is {accuracy:.2f}%, Average F1(weighted) is {f1_wai
78     print(f"TEST #{x}: Average Accuracy is {accuracy:.2f}%, Average TPR is {tpr:.2f}%, Avera
79     print("")
80
81

```

1 #Initial Models used:

```

2 dt3 = tree.DecisionTreeClassifier(max_depth=3,criterion='gini')
3 dt4 = tree.DecisionTreeClassifier(max_depth=4,criterion='gini')
4 dt5 = tree.DecisionTreeClassifier(max_depth=5,criterion='gini')
5 dt6 = tree.DecisionTreeClassifier(max_depth=6,criterion='gini')
6 dt7 = tree.DecisionTreeClassifier(max_depth=7,criterion='gini')
7 dt8 = tree.DecisionTreeClassifier(max_depth=8,criterion='gini')
8 dt9 = tree.DecisionTreeClassifier(max_depth=9,criterion='gini')
9 dt10 = tree.DecisionTreeClassifier(max_depth=10,criterion='gini')

```

```
10 dt11 = tree.DecisionTreeClassifier(max_depth=11,criterion='gini')
11 lr = LogisticRegression(max_iter=2000)
12 knn = KNeighborsClassifier()
13 svm = SVC()
14 mn_nb = MultinomialNB()
15 ga_nb = GaussianNB()
16 rf = RandomForestClassifier()
17 xgb = XGBClassifier()
```

## Initial Validation:

```
1 results = {}
2 test_num = 0
3 folds = 10
4 classCol = "Target"
5 df_list = [df,df_normalized,df_onehot,df_outliers_addressed,df_smote,df_normalized_smote,d
6
7 for test in ["DT_3","DT_4","DT_5","DT_6","DT_7","DT_8","DT_9","DT_10","DT_11","GA_NB","LR"
8     if test == "GA_NB":
9         model = ga_nb
10         test_name = "Gaussian Naive Bayes(GA-NB)"
11     elif test == "LR":
12         model = lr
13         test_name = "Logistic Regression(LR)"
14     elif test == "SVM":
15         model = svm
16         test_name = "Support Vector Machines(SVM)"
17     elif test == "KNN":
18         model = knn
19         test_name = "K Nearest Neighbours(KNN)"
20     elif test == "RF":
21         model = rf
22         test_name = "Random Forest(RF)"
23     elif test == "XGB":
24         model = xgb
25         test_name = "XG Boost(XGB)"
26     elif test == "DT_3":
27         model = dt3
28         test_name = "Decision Tree: 3 levels"
29     elif test == "DT_4":
30         model = dt4
31         test_name = "Decision Tree: 4 levels"
32     elif test == "DT_5":
33         model = dt5
34         test_name = "Decision Tree: 5 levels"
35     elif test == "DT_6":
36         model = dt6
37         test_name = "Decision Tree: 6 levels"
38     elif test == "DT_7":
```

```

39     model = dt7
40     test_name = "Decision Tree: 7 levels"
41 elif test == "DT_8":
42     model = dt8
43     test_name = "Decision Tree: 8 levels"
44 elif test == "DT_9":
45     model = dt9
46     test_name = "Decision Tree: 9 levels"
47 elif test == "DT_10":
48     model = dt10
49     test_name = "Decision Tree: 10 levels"
50 elif test == "DT_11":
51     model = dt11
52     test_name = "Decision Tree: 11 levels"
53
54 for idx,df_not_used in enumerate(df_list):
55     this_df = df_list[idx]
56     doCV(test_num,test_name,model,this_df,folds,classCol,1)
57     test_num += 1

```

For TEST #0: Decision Tree: 3 levels  
 For TEST #1: Decision Tree: 3 levels  
 For TEST #2: Decision Tree: 3 levels  
 For TEST #3: Decision Tree: 3 levels  
 For TEST #4: Decision Tree: 3 levels  
 For TEST #5: Decision Tree: 3 levels  
 For TEST #6: Decision Tree: 3 levels  
 For TEST #7: Decision Tree: 3 levels  
 For TEST #8: Decision Tree: 4 levels  
 For TEST #9: Decision Tree: 4 levels  
 For TEST #10: Decision Tree: 4 levels  
 For TEST #11: Decision Tree: 4 levels  
 For TEST #12: Decision Tree: 4 levels  
 For TEST #13: Decision Tree: 4 levels  
 For TEST #14: Decision Tree: 4 levels  
 For TEST #15: Decision Tree: 4 levels  
 For TEST #16: Decision Tree: 5 levels  
 For TEST #17: Decision Tree: 5 levels  
 For TEST #18: Decision Tree: 5 levels  
 For TEST #19: Decision Tree: 5 levels  
 For TEST #20: Decision Tree: 5 levels  
 For TEST #21: Decision Tree: 5 levels  
 For TEST #22: Decision Tree: 5 levels  
 For TEST #23: Decision Tree: 5 levels  
 For TEST #24: Decision Tree: 6 levels  
 For TEST #25: Decision Tree: 6 levels  
 For TEST #26: Decision Tree: 6 levels  
 For TEST #27: Decision Tree: 6 levels  
 For TEST #28: Decision Tree: 6 levels  
 For TEST #29: Decision Tree: 6 levels  
 For TEST #30: Decision Tree: 6 levels  
 For TEST #31: Decision Tree: 6 levels  
 For TEST #32: Decision Tree: 7 levels  
 For TEST #33: Decision Tree: 7 levels

After Cross-Val, Using Data-Set: Ori  
 After Cross-Val, Using Data-Set: Nor  
 After Cross-Val, Using Data-Set: Ori  
 After Cross-Val, Using Data-Set: Out  
 After Cross-Val, Using Data-Set: Ori  
 After Cross-Val, Using Data-Set: Nor  
 After Cross-Val, Using Data-Set: Ori  
 After Cross-Val, Using Data-Set: Out  
 After Cross-Val, Using Data-Set: Ori  
 After Cross-Val, Using Data-Set: Nor  
 After Cross-Val, Using Data-Set: Or  
 After Cross-Val, Using Data-Set: Ou  
 After Cross-Val, Using Data-Set: Or  
 After Cross-Val, Using Data-Set: No  
 After Cross-Val, Using Data-Set: Or  
 After Cross-Val, Using Data-Set: Ou  
 After Cross-Val, Using Data-Set: Or  
 After Cross-Val, Using Data-Set: No  
 After Cross-Val, Using Data-Set: Or  
 After Cross-Val, Using Data-Set: No  
 After Cross-Val, Using Data-Set: Or  
 After Cross-Val, Using Data-Set: Ou  
 After Cross-Val, Using Data-Set: Or  
 After Cross-Val, Using Data-Set: No  
 After Cross-Val, Using Data-Set: Or  
 After Cross-Val, Using Data-Set: Ou  
 After Cross-Val, Using Data-Set: Or  
 After Cross-Val, Using Data-Set: No  
 After Cross-Val, Using Data-Set: Or  
 After Cross-Val, Using Data-Set: No

For TEST #34: Decision Tree: 7 levels	After Cross-Val, Using Data-Set: Or
For TEST #35: Decision Tree: 7 levels	After Cross-Val, Using Data-Set: Ou
For TEST #36: Decision Tree: 7 levels	After Cross-Val, Using Data-Set: Or
For TEST #37: Decision Tree: 7 levels	After Cross-Val, Using Data-Set: No
For TEST #38: Decision Tree: 7 levels	After Cross-Val, Using Data-Set: Or
For TEST #39: Decision Tree: 7 levels	After Cross-Val, Using Data-Set: Ou
For TEST #40: Decision Tree: 8 levels	After Cross-Val, Using Data-Set: Or
For TEST #41: Decision Tree: 8 levels	After Cross-Val, Using Data-Set: No
For TEST #42: Decision Tree: 8 levels	After Cross-Val, Using Data-Set: Or
For TEST #43: Decision Tree: 8 levels	After Cross-Val, Using Data-Set: Ou
For TEST #44: Decision Tree: 8 levels	After Cross-Val, Using Data-Set: Or
For TEST #45: Decision Tree: 8 levels	After Cross-Val, Using Data-Set: No
For TEST #46: Decision Tree: 8 levels	After Cross-Val, Using Data-Set: Or
For TEST #47: Decision Tree: 8 levels	After Cross-Val, Using Data-Set: Ou
For TEST #48: Decision Tree: 9 levels	After Cross-Val, Using Data-Set: Or
For TEST #49: Decision Tree: 9 levels	After Cross-Val, Using Data-Set: No
For TEST #50: Decision Tree: 9 levels	After Cross-Val, Using Data-Set: Or
For TEST #51: Decision Tree: 9 levels	After Cross-Val, Using Data-Set: Ou
For TEST #52: Decision Tree: 9 levels	After Cross-Val, Using Data-Set: Or
For TEST #53: Decision Tree: 9 levels	After Cross-Val, Using Data-Set: No
For TEST #54: Decision Tree: 9 levels	After Cross-Val, Using Data-Set: Or
For TEST #55: Decision Tree: 9 levels	After Cross-Val, Using Data-Set: Ou
For TEST #56: Decision Tree: 10 levels	After Cross-Val, Using Data-Set: Or

```

1 #View these results again, this time just the top 10 using each sort method:
2 range_limit = min(10,results[f"LAST_TEST")+1) #Top 10 results desired.
3
4 results_list = list(range(0,results[f"LAST_TEST")+1))
5 results_list.sort(key=getTestAccuracy, reverse=True)
6
7 print("Results of ML Models: (sorted by accuracy) (top 10)")
8 print("")
9 for i in results_list[0:range_limit]: #Count starts at zero, so dont need to add one.
10  displayResult(i)
11 print("")
12 print("")
13
14
15 results_list = list(range(0,results[f"LAST_TEST")+1))
16 results_list.sort(key=getTestFNR, reverse=False)
17
18 print("Results of ML Models: (sorted by False Negative Rate(FNR)) (top 10)")
19 print("")
20 for i in results_list[0:range_limit]: #Count starts at zero, so dont need to add one.
21  displayResult(i)
22 print("")
23 print("")

```

Results of ML Models: (sorted by accuracy) (top 10)

Using ML Model: Random Forest(RF) and Original Data-Set, After SMOTE:

TEST #108: Average Accuracy is 87.57%, Average TPR is 87.62%, Average FNR is 12.38%



Using ML Model: XG Boost(XGB) and Original Data-Set, After SMOTE:

TEST #116: Average Accuracy is 86.87%, Average TPR is 86.93%, Average FNR is 13.07%

Using ML Model: Random Forest(RF) and Original Data-Set, After ONEHOT, After SMOTE:

TEST #110: Average Accuracy is 86.68%, Average TPR is 86.74%, Average FNR is 13.26%

Using ML Model: Random Forest(RF) and Outliers Addressed Data-Set, After ONEHOT, After

TEST #111: Average Accuracy is 86.48%, Average TPR is 86.53%, Average FNR is 13.47%

Using ML Model: XG Boost(XGB) and Original Data-Set, After ONEHOT, After SMOTE:

TEST #118: Average Accuracy is 86.48%, Average TPR is 86.53%, Average FNR is 13.47%

Using ML Model: Random Forest(RF) and Normalized Data-Set, After SMOTE:

TEST #109: Average Accuracy is 86.09%, Average TPR is 86.14%, Average FNR is 13.86%

Using ML Model: XG Boost(XGB) and Normalized Data-Set, After SMOTE:

TEST #117: Average Accuracy is 85.98%, Average TPR is 86.03%, Average FNR is 13.97%

Using ML Model: Random Forest(RF) and Outliers Addressed Data-Set:

TEST #107: Average Accuracy is 85.91%, Average TPR is 85.64%, Average FNR is 14.36%

Using ML Model: Random Forest(RF) and Normalized Data-Set:

TEST #105: Average Accuracy is 85.91%, Average TPR is 85.54%, Average FNR is 14.46%

Using ML Model: Logistic Regression(LR) and Outliers Addressed Data-Set, After ONEHOT

TEST #87: Average Accuracy is 85.79%, Average TPR is 85.85%, Average FNR is 14.15%

Results of ML Models: (sorted by False Negative Rate(FNR)) (top 10)

Using ML Model: Random Forest(RF) and Original Data-Set, After SMOTE:

TEST #108: Average Accuracy is 87.57%, Average TPR is 87.62%, Average FNR is 12.38%

Using ML Model: XG Boost(XGB) and Original Data-Set, After SMOTE:

TEST #116: Average Accuracy is 86.87%, Average TPR is 86.93%, Average FNR is 13.07%

Using ML Model: Random Forest(RF) and Original Data-Set, After ONEHOT, After SMOTE:

TEST #110: Average Accuracy is 86.68%, Average TPR is 86.74%, Average FNR is 13.26%

Using ML Model: Random Forest(RF) and Outliers Addressed Data-Set, After ONEHOT, After

TEST #111: Average Accuracy is 86.48%, Average TPR is 86.53%, Average FNR is 13.47%

Using ML Model: XG Boost(XGB) and Original Data-Set, After ONEHOT, After SMOTE:

TEST #118: Average Accuracy is 86.48%, Average TPR is 86.53%, Average FNR is 13.47%

Using ML Model: Random Forest(RF) and Normalized Data-Set, After SMOTE:

TEST #109: Average Accuracy is 86.09%, Average TPR is 86.14%, Average FNR is 13.86%

Using ML Model: XG Boost(XGB) and Normalized Data-Set, After SMOTE:

TEST #117: Average Accuracy is 85.98%, Average TPR is 86.03%, Average FNR is 13.97%

## Hypertuning:

```

1 #For Hyper-tuning and secondary validation, only three dataset variations will be used.
2 #"Original Data-Set, After ONEHOT, After SMOTE"
3 #"Outliers Addressed Data-Set, After ONEHOT, After SMOTE"
4 #"Original Data-Set, After SMOTE"
5 #
6 #Copying these over to easier variables for these two stages.
7 df1 = df_smote.copy()
8 df1.name = df_smote.name #For some reason df.copy() doesn't copy the "name" attribute.
9
10 df2 = df_onehot_smote.copy()
11 df2.name = df_onehot_smote.name #For some reason df.copy() doesn't copy the "name" attribu
12
13 df3 = df_onehot_outliers_addressed_smote.copy()
14 df3.name = df_onehot_outliers_addressed_smote.name #For some reason df.copy() doesn't copy
15
16 df_list = []
17 df_list = [df1, df2, df3]
18
19 classCol = "Target"

```

```

1 #Based on the initial results, for hyper-tuning, focusing on models RF, XGB, and LR.
2
3 #Since hyper-tuning takes the most time to run, of everything in this project,
4 #wrapping this section up in a "if" statement so I can turn it off for future top down run
5 #Once hypertuned values are retrieved, the output will be copied into a text block for ref
6 #perform_tuning = True
7 perform_tuning = False
8
9 #Models used for Hyper-tuning.
10 rf = RandomForestClassifier()
11 xgb = XGBClassifier()
12 lr = LogisticRegression(max_iter=10000)
13
14 #Parameter Options tried for Hyper-tuning.
15 params_rf = {'n_estimators':[100,200,300,400,500], 'min_samples_leaf':[5, 10, 15, 20, 25,
16 params_xgb = {'n_estimators': [100,200,300,400,500,600,700,800,900,1000], 'learning_rate':
17 params_lr = {'solver':['newton-cg', 'lbfgs', 'sag', 'saga', 'liblinear'], 'penalty':['l2']}
18
19 folds = 20
20
21 if perform_tuning:
22     params_rf = {'n_estimators':[100,200,300,400,500], 'min_samples_leaf':[5, 10, 15, 20, 25
23     params_xgb = {'n_estimators': [100,200,300,400,500,600,700,800,900,1000], 'learning_rate
24     params_lr = {'solver':['newton-cg', 'lbfgs', 'sag', 'saga', 'liblinear'], 'penalty':['l2
25
26     for idx,df_not_used in enumerate(df_list):
27         this_df = df_list[idx]
28
29         feature_names=this_df.columns[this_df.columns != classCol]

```

```

30 #80% training and 20% test
31 X_train, X_test, Y_train, Y_test = train_test_split(this_df.loc[:, feature_names], thi
32
33 grid_rf = GridSearchCV(rf, param_grid=params_rf, cv=folds)
34 grid_rf.fit(X_train, Y_train)
35 print(f"Using Data-Set: {this_df.name}:")
36 print("Hyper-Tuned Parameters for Random Forest:", grid_rf.best_params_)
37 print("")
38
39 rs_xgb = RandomizedSearchCV(xgb, param_distributions=params_xgb, cv=folds)
40 rs_xgb.fit(X_train, Y_train)
41 print(f"Using Data-Set: {this_df.name}:")
42 print("Hyper-Tuned Parameters for XGBoost:", rs_xgb.best_params_)
43 print("")
44
45 grid_lr = GridSearchCV(lr, param_grid=params_lr, cv=folds)
46 grid_lr.fit(X_train, Y_train)
47 print(f"Using Data-Set: {this_df.name}:")
48 print("Hyper-Tuned Parameters for Logistic Regression:", grid_lr.best_params_)
49 print("")
50

```

### Results from hyper-tuning:

Using Data-Set: Original Data-Set, After SMOTE: Hyper-Tuned Parameters for Random Forest:  
{'min\_samples\_leaf': 10, 'n\_estimators': 500}

Using Data-Set: Original Data-Set, After SMOTE: Hyper-Tuned Parameters for XGBoost:  
{'n\_estimators': 1000, 'learning\_rate': 0.3}

Using Data-Set: Original Data-Set, After SMOTE: Hyper-Tuned Parameters for Logistic Regression:  
{'C': 100, 'penalty': 'l2', 'solver': 'newton-cg'}

Using Data-Set: Original Data-Set, After ONEHOT, After SMOTE: Hyper-Tuned Parameters for  
Random Forest: {'min\_samples\_leaf': 5, 'n\_estimators': 400}

Using Data-Set: Original Data-Set, After ONEHOT, After SMOTE: Hyper-Tuned Parameters for  
XGBoost: {'n\_estimators': 500, 'learning\_rate': 0.2}

Using Data-Set: Original Data-Set, After ONEHOT, After SMOTE: Hyper-Tuned Parameters for  
Logistic Regression: {'C': 100, 'penalty': 'l2', 'solver': 'newton-cg'}

Using Data-Set: Outliers Addressed Data-Set, After ONEHOT, After SMOTE: Hyper-Tuned Parameters  
for Random Forest: {'min\_samples\_leaf': 5, 'n\_estimators': 200}

Using Data-Set: Outliers Addressed Data-Set, After ONEHOT, After SMOTE: Hyper-Tuned Parameters  
for XGBoost: {'n\_estimators': 400, 'learning\_rate': 0.5}

## Using Data-Set: Outliers Addressed Data-Set, After ONEHOT, After SMOTE: Hyper-Tuned Parameters

for Logistic Regression: {'C': 10, 'penalty': 'l2', 'solver': 'liblinear'}

```

1 #Apply Hyper-tuned parameters.
2 #I know there is an automated way of applying the Hyper-tuned parameters from the output o
3 #grid search, but this works too.
4
5 rf_tuned_df1 = RandomForestClassifier(min_samples_leaf = 10, n_estimators = 500)
6 xgb_tuned_df1 = XGBClassifier(n_estimators = 1000, learning_rate = 0.3)
7 lr_tuned_df1 = LogisticRegression(max_iter=10000, C=100, penalty="l2", solver="newton-cg")
8
9 rf_tuned_df2 = RandomForestClassifier(min_samples_leaf = 5, n_estimators = 400)
10 xgb_tuned_df2 = XGBClassifier(n_estimators = 500, learning_rate = 0.2)
11 lr_tuned_df2 = LogisticRegression(max_iter=10000, C=100, penalty="l2", solver="newton-cg")
12
13 rf_tuned_df3 = RandomForestClassifier(min_samples_leaf = 5, n_estimators = 200)
14 xgb_tuned_df3 = XGBClassifier(n_estimators = 400, learning_rate = 0.5)
15 lr_tuned_df3 = LogisticRegression(max_iter=10000, C=10, penalty="l2", solver="liblinear")
16

```

## Secondary Validation: (After Hypertuning):

```

1 #This will start testing all over again.
2 results = {}
3 test_num = 0
4 folds = 20
5
6 for test in ["LR_DF1", "LR_DF2", "LR_DF3", "LR_TUNED_DF1", "LR_TUNED_DF2", "LR_TUNED_DF3",
7             "RF_DF1", "RF_DF2", "RF_DF3", "RF_TUNED_DF1", "RF_TUNED_DF2", "RF_TUNED_DF3",
8             "XGB_DF1", "XGB_DF2", "XGB_DF3", "XGB_TUNED_DF1", "XGB_TUNED_DF2", "XGB_TUNED_DF3"]
9     if test == "LR_DF1":
10         model = lr
11         test_name = "Logistic Regression(LR)"
12         this_df = df1
13     elif test == "LR_DF2":
14         model = lr
15         test_name = "Logistic Regression(LR)"
16         this_df = df2
17     elif test == "LR_DF3":
18         model = lr
19         test_name = "Logistic Regression(LR)"
20         this_df = df3
21
22     elif test == "LR_TUNED_DF1":
23         model = lr_tuned_df1
24         test_name = "Logistic Regression(LR) - Tuned"
25         this_df = df1
26     elif test == "LR_TUNED_DF2":
27         model = lr_tuned_df2

```

```
28     test_name = "Logistic Regression(LR) - Tuned"
29     this_df = df2
30 elif test == "LR_TUNED_DF3":
31     model = lr_tuned_df3
32     test_name = "Logistic Regression(LR) - Tuned"
33     this_df = df3
34
35 elif test == "RF_DF1":
36     model = rf
37     test_name = "Random Forest(RF)"
38     this_df = df1
39 elif test == "RF_DF2":
40     model = rf
41     test_name = "Random Forest(RF)"
42     this_df = df2
43 elif test == "RF_DF3":
44     model = rf
45     test_name = "Random Forest(RF)"
46     this_df = df3
47
48 elif test == "RF_TUNED_DF1":
49     model = rf_tuned_df1
50     test_name = "Random Forest(RF) - Tuned"
51     this_df = df1
52 elif test == "RF_TUNED_DF2":
53     model = rf_tuned_df2
54     test_name = "Random Forest(RF) - Tuned"
55     this_df = df2
56 elif test == "RF_TUNED_DF3":
57     model = rf_tuned_df3
58     test_name = "Random Forest(RF) - Tuned"
59     this_df = df3
60
61 elif test == "XGB_DF1":
62     model = xgb
63     test_name = "XG Boost(XGB)"
64     this_df = df1
65 elif test == "XGB_DF2":
66     model = xgb
67     test_name = "XG Boost(XGB)"
68     this_df = df2
69 elif test == "XGB_DF3":
70     model = xgb
71     test_name = "XG Boost(XGB)"
72     this_df = df3
73
74 elif test == "XGB_TUNED_DF1":
75     model = xgb_tuned_df1
76     test_name = "XG Boost(XGB) - Tuned"
77     this_df = df1
78 elif test == "XGB_TUNED_DF2":
```

```

79     model = xgb_tuned_df2
80     test_name = "XG Boost(XGB) - Tuned"
81     this_df = df2
82 elif test == "XGB_TUNED_DF3":
83     model = xgb_tuned_df3
84     test_name = "XG Boost(XGB) - Tuned"
85     this_df = df3
86
87 doCV(test_num,test_name,model,this_df,folds,classCol,1)
88 test_num += 1

```

For TEST #0: Logistic Regression(LR)	After Cross-Val, Using Data-Set: Original
For TEST #1: Logistic Regression(LR)	After Cross-Val, Using Data-Set: Original
For TEST #2: Logistic Regression(LR)	After Cross-Val, Using Data-Set: Outliers
For TEST #3: Logistic Regression(LR) - Tuned	After Cross-Val, Using Data-Set: Original
For TEST #4: Logistic Regression(LR) - Tuned	After Cross-Val, Using Data-Set: Original
For TEST #5: Logistic Regression(LR) - Tuned	After Cross-Val, Using Data-Set: Outliers
For TEST #6: Random Forest(RF)	After Cross-Val, Using Data-Set: Original
For TEST #7: Random Forest(RF)	After Cross-Val, Using Data-Set: Original
For TEST #8: Random Forest(RF)	After Cross-Val, Using Data-Set: Outliers
For TEST #9: Random Forest(RF) - Tuned	After Cross-Val, Using Data-Set: Original
For TEST #10: Random Forest(RF) - Tuned	After Cross-Val, Using Data-Set: Original
For TEST #11: Random Forest(RF) - Tuned	After Cross-Val, Using Data-Set: Outliers
For TEST #12: XG Boost(XGB)	After Cross-Val, Using Data-Set: Original
For TEST #13: XG Boost(XGB)	After Cross-Val, Using Data-Set: Original
For TEST #14: XG Boost(XGB)	After Cross-Val, Using Data-Set: Outliers
For TEST #15: XG Boost(XGB) - Tuned	After Cross-Val, Using Data-Set: Original
For TEST #16: XG Boost(XGB) - Tuned	After Cross-Val, Using Data-Set: Original
For TEST #17: XG Boost(XGB) - Tuned	After Cross-Val, Using Data-Set: Outliers

```

1 #View these results, after hyper-tuning, just the top 5 using each sort method:
2 range_limit = min(5,results[f"LAST_TEST"]+1) #Top 5 results desired.
3
4 results_list = list(range(0,results[f"LAST_TEST"]+1))
5 results_list.sort(key=getTestAccuracy, reverse=True)
6
7 print("Results of ML Models: (sorted by accuracy) (top 5)")
8 print("")
9 for i in results_list[0:range_limit]: #Count starts at zero, so dont need to add one.
10     displayResult(i)
11 print("")
12 print("")
13
14
15 results_list = list(range(0,results[f"LAST_TEST"]+1))
16 results_list.sort(key=getTestFNR, reverse=False)
17
18 print("Results of ML Models: (sorted by False Negative Rate(FNR)) (top 5)")
19 print("")
20 for i in results_list[0:range_limit]: #Count starts at zero, so dont need to add one.
21     displayResult(i)

```

```
22 print("")
23 print("")
24
25
26 #View these results, after hyper-tuning, just the bottom 5 using each sort method:
27 range_limit = min(5,results[f"LAST_TEST")+1) #Bottom 5 results desired.
28
29 results_list = list(range(0,results[f"LAST_TEST")+1))
30 results_list.sort(key=getTestAccuracy, reverse=False)
31
32 print("Results of ML Models: (sorted by accuracy) (bottom 5)")
33 print("")
34 for i in results_list[0:range_limit]: #Count starts at zero, so dont need to add one.
35     displayResult(i)
36 print("")
37 print("")
38
39
```

Results of ML Models: (sorted by accuracy) (top 5)

Using ML Model: Random Forest(RF) and Original Data-Set, After SMOTE:

TEST #6: Average Accuracy is 88.44%, Average TPR is 88.37%, Average FNR is 11.63%

Using ML Model: Random Forest(RF) - Tuned and Original Data-Set, After ONEHOT, After

TEST #10: Average Accuracy is 88.24%, Average TPR is 88.18%, Average FNR is 11.82%

Using ML Model: XG Boost(XGB) and Original Data-Set, After SMOTE:

TEST #12: Average Accuracy is 88.24%, Average TPR is 88.17%, Average FNR is 11.83%

Using ML Model: Random Forest(RF) and Original Data-Set, After ONEHOT, After SMOTE:

TEST #7: Average Accuracy is 88.14%, Average TPR is 88.09%, Average FNR is 11.91%

Using ML Model: Random Forest(RF) and Outliers Addressed Data-Set, After ONEHOT, After

TEST #8: Average Accuracy is 87.85%, Average TPR is 87.78%, Average FNR is 12.22%

Results of ML Models: (sorted by False Negative Rate(FNR)) (top 5)

Using ML Model: Random Forest(RF) and Original Data-Set, After SMOTE:

TEST #6: Average Accuracy is 88.44%, Average TPR is 88.37%, Average FNR is 11.63%

Using ML Model: Random Forest(RF) - Tuned and Original Data-Set, After ONEHOT, After

TEST #10: Average Accuracy is 88.24%, Average TPR is 88.18%, Average FNR is 11.82%

Using ML Model: XG Boost(XGB) and Original Data-Set, After SMOTE:

TEST #12: Average Accuracy is 88.24%, Average TPR is 88.17%, Average FNR is 11.83%

Using ML Model: Random Forest(RF) and Original Data-Set, After ONEHOT, After SMOTE:

TEST #7: Average Accuracy is 88.14%, Average TPR is 88.09%, Average FNR is 11.91%

Using ML Model: Random Forest(RF) and Outliers Addressed Data-Set, After ONEHOT, After

TEST #8: Average Accuracy is 87.85%, Average TPR is 87.78%, Average FNR is 12.22%

Results of ML Models: (sorted by accuracy) (bottom 5)

Using ML Model: Logistic Regression(LR) - Tuned and Original Data-Set, After SMOTE:  
TEST #3: Average Accuracy is 85.30%, Average TPR is 85.19%, Average FNR is 14.81%

Using ML Model: Logistic Regression(LR) and Original Data-Set, After SMOTE:  
TEST #0: Average Accuracy is 85.30%, Average TPR is 85.19%, Average FNR is 14.81%

Using ML Model: XG Boost(XGB) - Tuned and Outliers Addressed Data-Set, After ONEHOT, ,  
TEST #17: Average Accuracy is 86.17%, Average TPR is 86.10%, Average FNR is 13.90%

Using ML Model: XG Boost(XGB) - Tuned and Original Data-Set, After SMOTE:  
TEST #15: Average Accuracy is 86.36%, Average TPR is 86.28%, Average FNR is 13.72%

Using ML Model: Logistic Regression(LR) - Tuned and Original Data-Set, After ONEHOT, ,  
TEST #4: Average Accuracy is 86.37%, Average TPR is 86.28%, Average FNR is 13.72%

```

1 #Comparing Results.
2 #Due to random splits, the results change from run to run, but for the most part XGBoost w
3 #whereas Random Forest produces very similar results tuned, or not tuned. Consistently th
4 #
5 #
6 #Using ML Model: Random Forest(RF) and Original Data-Set, After SMOTE:
7 #TEST #6: Average Accuracy is 88.44%, Average TPR is 88.37%, Average FNR is 11.63%
8
9 #Using ML Model: Random Forest(RF) - Tuned and Original Data-Set, After ONEHOT, After SMOT
10 #TEST #10: Average Accuracy is 88.24%, Average TPR is 88.18%, Average FNR is 11.82%
11
12 #Using ML Model: XG Boost(XGB) and Original Data-Set, After SMOTE:
13 #TEST #12: Average Accuracy is 88.24%, Average TPR is 88.17%, Average FNR is 11.83%
14
15 #Using ML Model: Random Forest(RF) and Original Data-Set, After ONEHOT, After SMOTE:
16 #TEST #7: Average Accuracy is 88.14%, Average TPR is 88.09%, Average FNR is 11.91%
17
18 #
19 # But are the differences between models, and data-sets statistically significant?
20 #
21 # For example, we have these three pairs of results:
22 #
23 # Pair #1: Same Data-set/Different models.
24 #Using ML Model: Random Forest(RF) and Original Data-Set, After SMOTE:
25 #TEST #6: Average Accuracy is 88.44%, Average TPR is 88.37%, Average FNR is 11.63%
26 #
27 #Using ML Model: XG Boost(XGB) and Original Data-Set, After SMOTE:
28 #TEST #12: Average Accuracy is 88.24%, Average TPR is 88.17%, Average FNR is 11.83%
29 #

```



```

30 #
31 # Pair #2: Same Data-set/RF vs RF Tuned
32 #Using ML Model: Random Forest(RF) - Tuned and Original Data-Set, After ONEHOT, After SMOT
33 #TEST #10: Average Accuracy is 88.24%, Average TPR is 88.18%, Average FNR is 11.82%
34
35 #Using ML Model: Random Forest(RF) and Original Data-Set, After ONEHOT, After SMOTE:
36 #TEST #7: Average Accuracy is 88.14%, Average TPR is 88.09%, Average FNR is 11.91%
37 #
38 # Each pair uses the same data-set, which make them suitable to compare results against
39 # each other using a paired dependant t-test.
40 #
41 #
42 def my_ttest(x,y,verbose=0):
43     t,p = stats.ttest_rel(getTestAccVec(x),getTestAccVec(y))
44     x_name = results[f"TEST{x}#NAME"]
45     x_df_name = results[f"TEST{x}#DFNAME"]
46     y_name = results[f"TEST{y}#NAME"]
47     y_df_name = results[f"TEST{y}#DFNAME"]
48
49     if p >= 0.05:
50         if verbose > 0:
51             print(f"TEST #{x}: {x_name}: Using Data-Set: {x_df_name} &")
52             print(f"TEST #{y}: {y_name}: Using Data-Set: {y_df_name}")
53             print(f"Results for Accuracy are NOT significantly different from each other, with t
54             print("")
55         return(False)
56     else:
57         if verbose > 0:
58             print(f"TEST #{x}: {x_name}: Using Data-Set: {x_df_name} &")
59             print(f"TEST #{y}: {y_name}: Using Data-Set: {y_df_name}")
60             print(f"Results for Accuracy are significantly different from each other, with t={t}
61             print("")
62         return(True)
63
64 #Pair #1:
65 res = my_ttest(6,12,1)
66
67 #Pair #2:
68 res = my_ttest(10,7,1)

```

TEST #6: Random Forest(RF): Using Data-Set: Original Data-Set, After SMOTE &  
 TEST #12: XG Boost(XGB): Using Data-Set: Original Data-Set, After SMOTE  
 Results for Accuracy are NOT significantly different from each other, with t=0.261847487

TEST #10: Random Forest(RF) - Tuned: Using Data-Set: Original Data-Set, After ONEHOT, A  
 TEST #7: Random Forest(RF): Using Data-Set: Original Data-Set, After ONEHOT, After SMOTE  
 Results for Accuracy are NOT significantly different from each other, with t=0.197923169



1 #The paired dependant t-test above revealed there is no significant difference between any